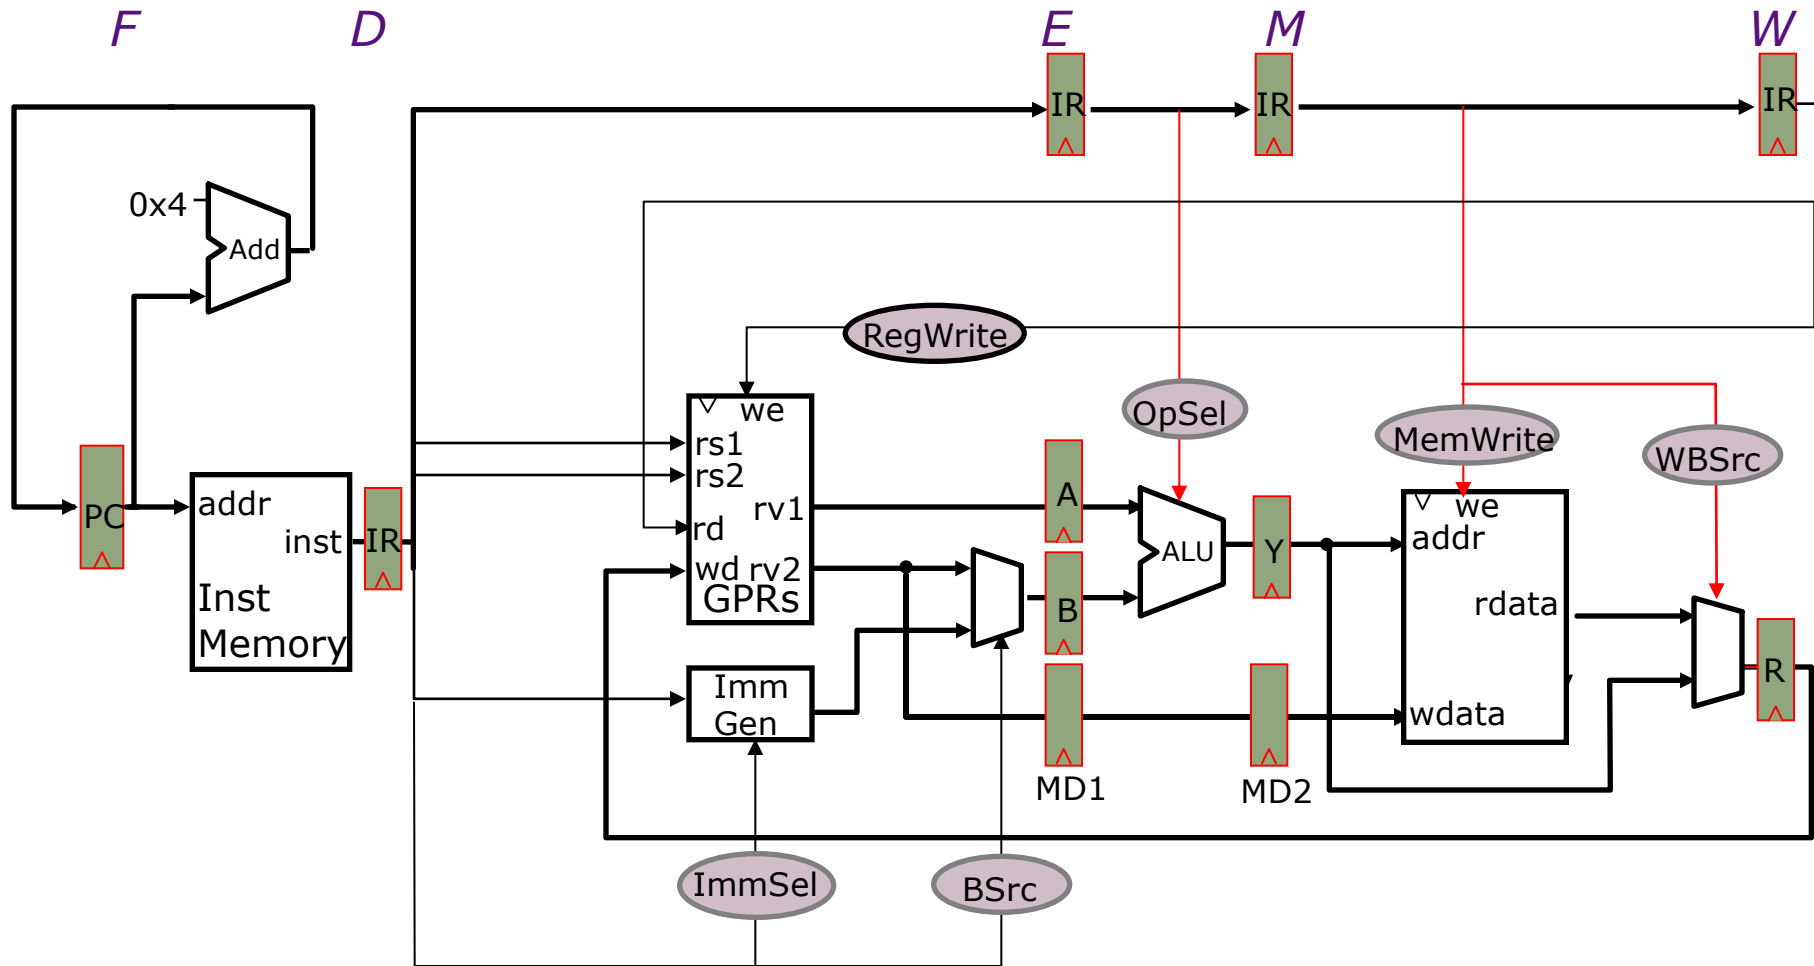


Instruction Pipelining: Hazard Resolution, Timing Constraints

Daniel Sanchez

Computer Science and Artificial Intelligence Laboratory
M.I.T.

Reminder: Pipelined RISC-V Datapath *without jumps*



Pipeline Diagram – Ideal Pipelining

<i>time</i>	t0	t1	t2	t3	t4	t5	t6	t7	...
(I ₁) $a1 \leftarrow a0 + 10$	IF ₁	ID ₁	EX ₁	MA ₁	WB ₁				
(I ₂) $a3 \leftarrow a2 + 12$		IF ₂	ID ₂	EX ₂	MA ₂	WB ₂			
(I ₃) $a5 \leftarrow a4 + 14$			IF ₃	ID ₃	EX ₃	MA ₃	WB ₃		
(I ₄) $a7 \leftarrow a6 + 16$				IF ₄	ID ₄	EX ₄	MA ₄	WB ₄	
(I ₅) $s1 \leftarrow s0 + 18$					IF ₅	ID ₅	EX ₅	MA ₅	WB ₅

In steady state, what are the ...

CPI?

IPC?

Cycles to exec each instr?

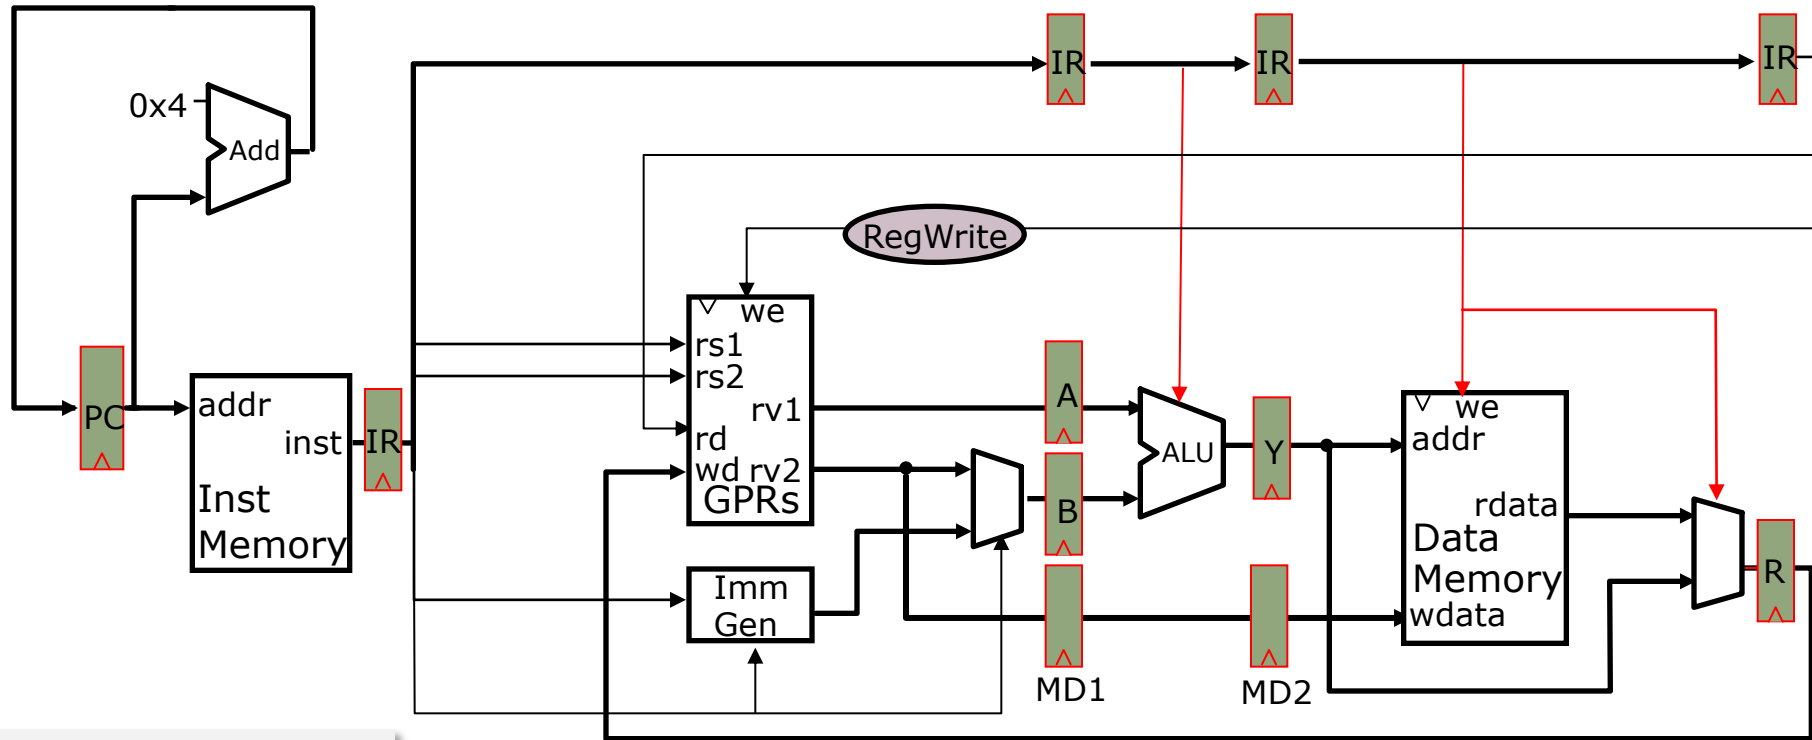
Num instrs in flight?

Pipelining increases clock frequency,
but instruction dependences may increase CPI

How instructions can interact with each other in a pipeline

- An instruction in the pipeline may need a resource being used by another instruction in the pipeline
→ *structural hazard*
- An instruction may depend on a value produced by an earlier instruction
 - Dependence may be for a data calculation
→ *data hazard*
 - Dependence may be for calculating the next PC
→ *control hazard (branches, interrupts)*

Data Hazards



```
...  
a1 ← a0 + 10  
a3 ← a1 + 12  
...
```

Pipeline Diagram – Hazard

<i>time</i>	t0	t1	t2	t3	t4	t5	t6	t7
(I ₁) $a1 \leftarrow a0 + 10$	IF ₁	ID ₁	EX ₁	MA ₁	WB ₁				
(I ₂) $a3 \leftarrow a1 + 12$		IF ₂	ID ₂	EX ₂	MA ₂	WB ₂			

Resolving Data Hazards

Strategy 1: *Wait for the result to be available by freezing earlier pipeline stages → **stall***

Strategy 2: *Route data as soon as possible after it is calculated to the earlier pipeline stage → **bypass***

Strategy 3: ***Speculate*** on the dependence
Two cases:

Guessed correctly → no special action required
Guessed incorrectly → kill and restart

Resolving Data Hazards

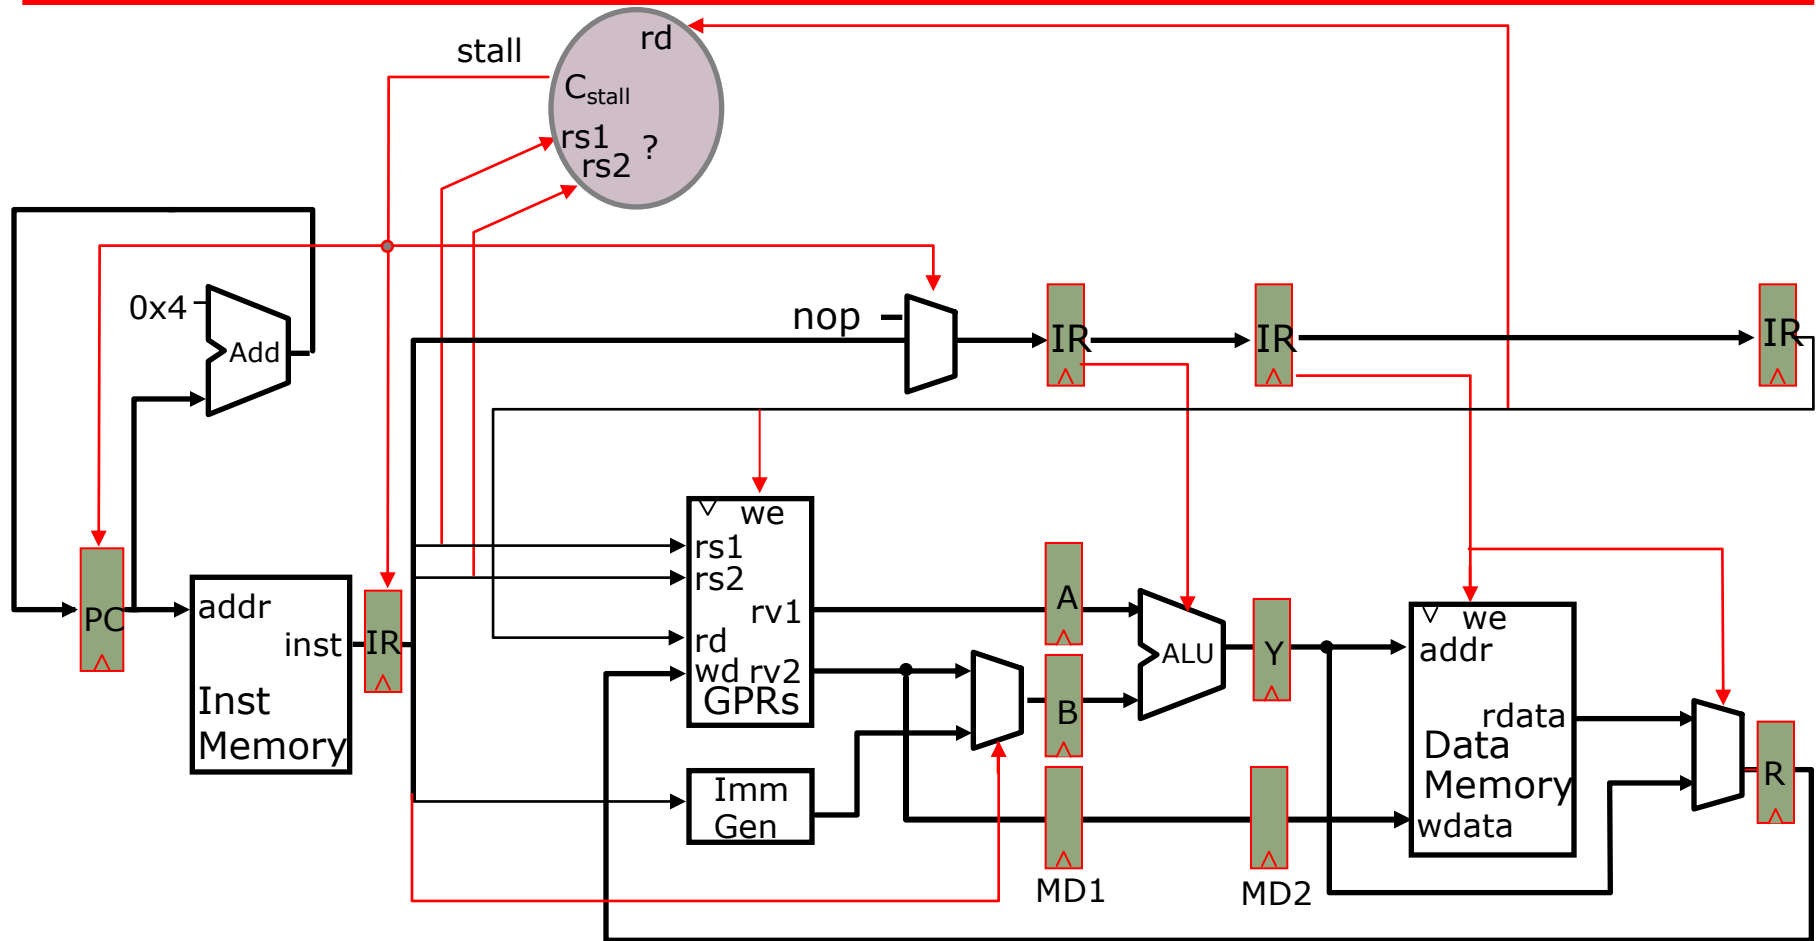
Strategy 1: *Wait for the result to be available by freezing earlier pipeline stages → **stall***



Stall DEC & IF when instruction in DEC reads a register that is written by any earlier in-flight instruction (in EXE, MEM, or WB)

Reminder: Stall Control Logic

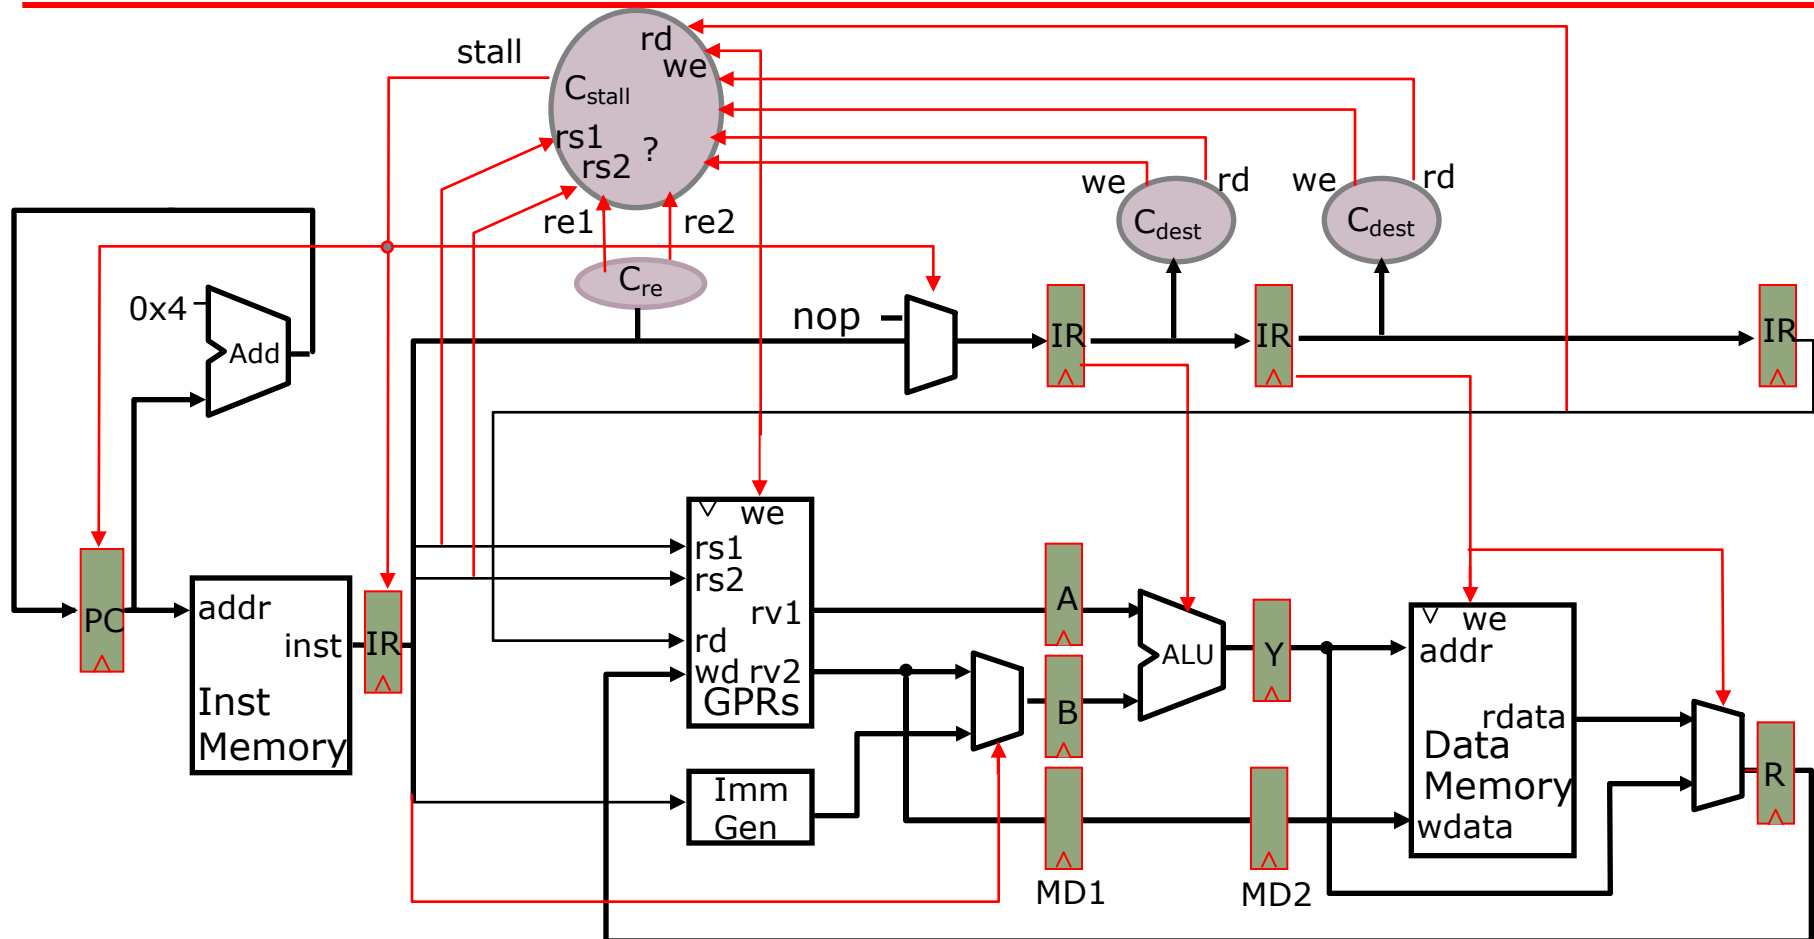
ignoring jumps & branches



Compare the *source registers* of the instruction in the decode stage with the *destination register* of the *uncommitted instructions*.

Reminder: Stall Control Logic

ignoring jumps & branches



Should we always stall if the rs field(s) matches some rd?

Source & Destination Registers

R-type:

funct7	rs2	rs1	funct3	rd	opcode
--------	-----	-----	--------	----	--------

I-type:

imm[11:0]	rs1	funct3	rd	opcode	
imm[11:5]	rs2	rs1	funct3	imm[4:0]	opcode

		<i>source(s)</i>	<i>dest</i>
ALU	$rd \leftarrow (rs1) \text{ func } (rs2)$	rs1, rs2	rd
ALUi	$rd \leftarrow (rs1) \text{ func } \text{imm}$	rs1	rd
LW	$rd \leftarrow M[(rs1) + \text{imm}]$	rs1	rd
SW	$M[(rs1) + \text{imm}] \leftarrow (rs2)$	rs1, rs2	
BRANCH	$\text{cond } (rs1) \text{ func } (rs2)$		
	<i>true:</i> $PC \leftarrow (PC) + \text{imm}$	rs1, rs2	
	<i>false:</i> $PC \leftarrow (PC) + 4$	rs1, rs2	
JAL	$rd \leftarrow (PC+4), PC \leftarrow (PC) + \text{imm}$		rd
JALR	$rd \leftarrow (PC+4), PC \leftarrow (rs1) + \text{imm}^*$	rs1	rd

**More precisely: $pc \leftarrow \{(reg[rs1] + imm)[31:1], 1'b0\}$*

Deriving the Stall Signal

C_{dest}

we = Case opcode
ALU, ALUi, LW, JAL, JALR
...
 $\Rightarrow (rd \neq 0)$
 \Rightarrow off

C_{re}

re1 = Case opcode
ALU, ALUi,
LW, SW, BRANCH,
JR, JALR \Rightarrow on

re2 = Case opcode
ALU, SW, BRANCH \Rightarrow on
... \Rightarrow off

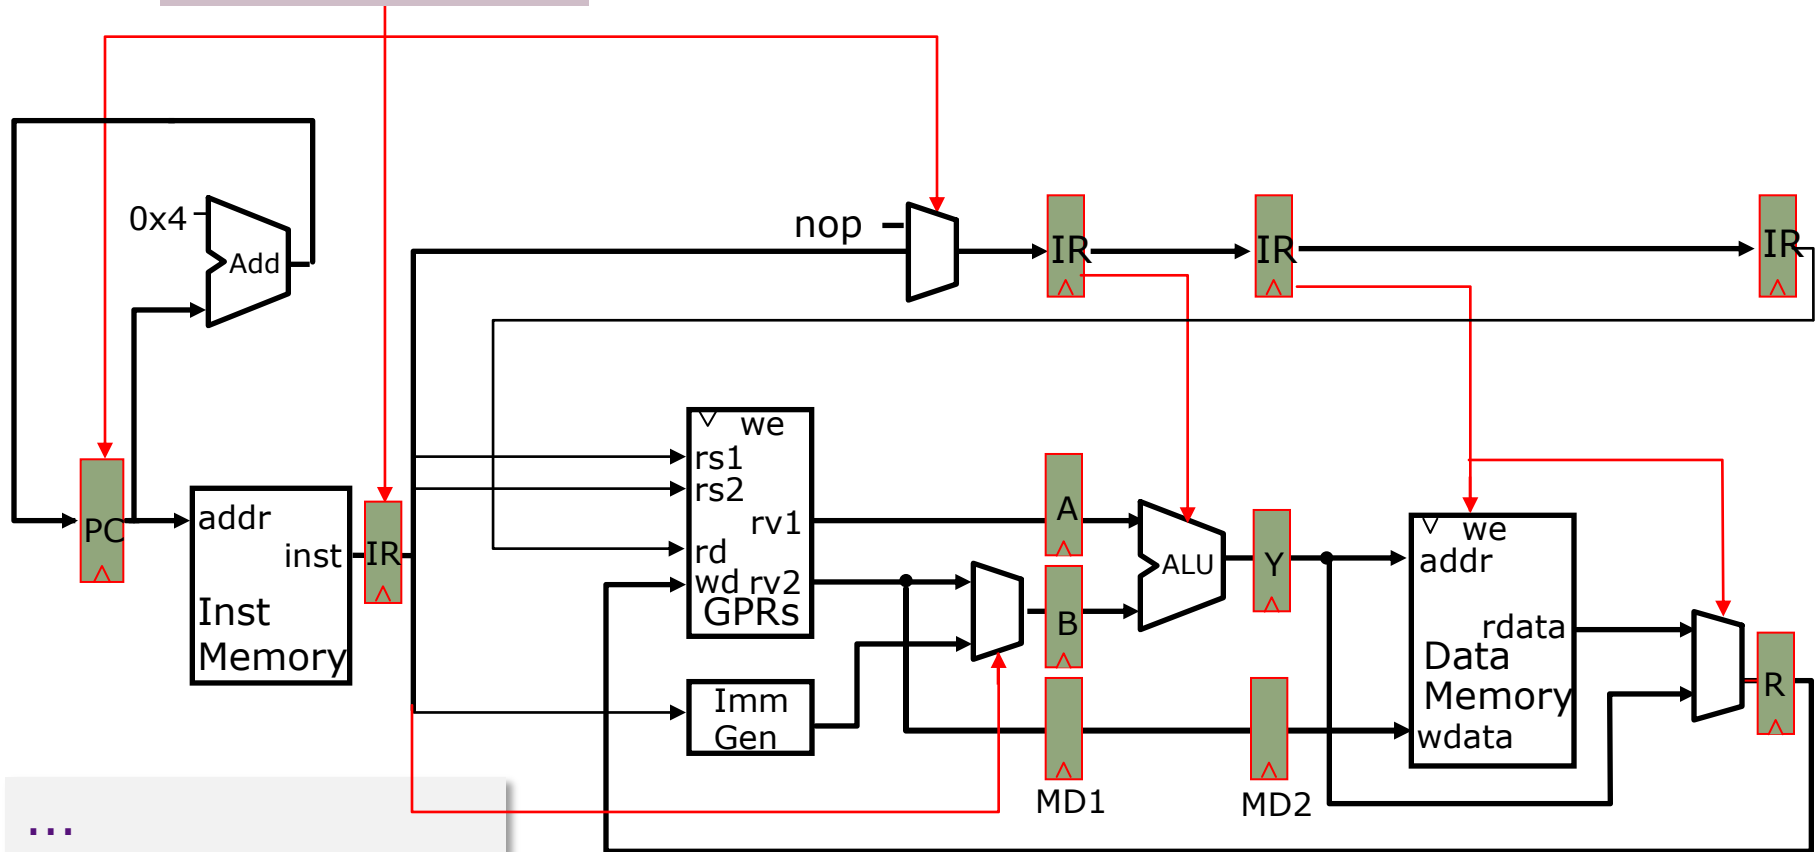
C_{stall}

stall = $((rs1_D == rd_E) \cdot we_E +$
 $(rs1_D == rd_M) \cdot we_M +$
 $(rs1_D == rd_W) \cdot we_W) \cdot re1_D$
+
 $((rs2_D == rd_E) \cdot we_E +$
 $(rs2_D == rd_M) \cdot we_M +$
 $(rs2_D == rd_W) \cdot we_W) \cdot re2_D$

This is not
the full story !

Hazards due to Loads & Stores

Stall Condition



...

$M[(a1)+7] \leftarrow (a2)$

$a4 \leftarrow M[(a3)+5]$

...

Is there any possible data hazard in this instruction sequence?

Load & Store Hazards

```
...  
M[(a1)+7] ← (a2)  
a4 ← M[(a3)+5]  
...
```

$(a1)+7 = (a3)+5 \Rightarrow \text{data hazard}$

However, the hazard is avoided because *our memory system completes writes in one cycle!*

Load/Store hazards are sometimes resolved in the pipeline and sometimes in the memory system itself.


More on this later in the course.

Resolving Data Hazards (2)

Strategy 2:


Route data as soon as possible after it is calculated to the earlier pipeline stage → *bypass*

Bypassing



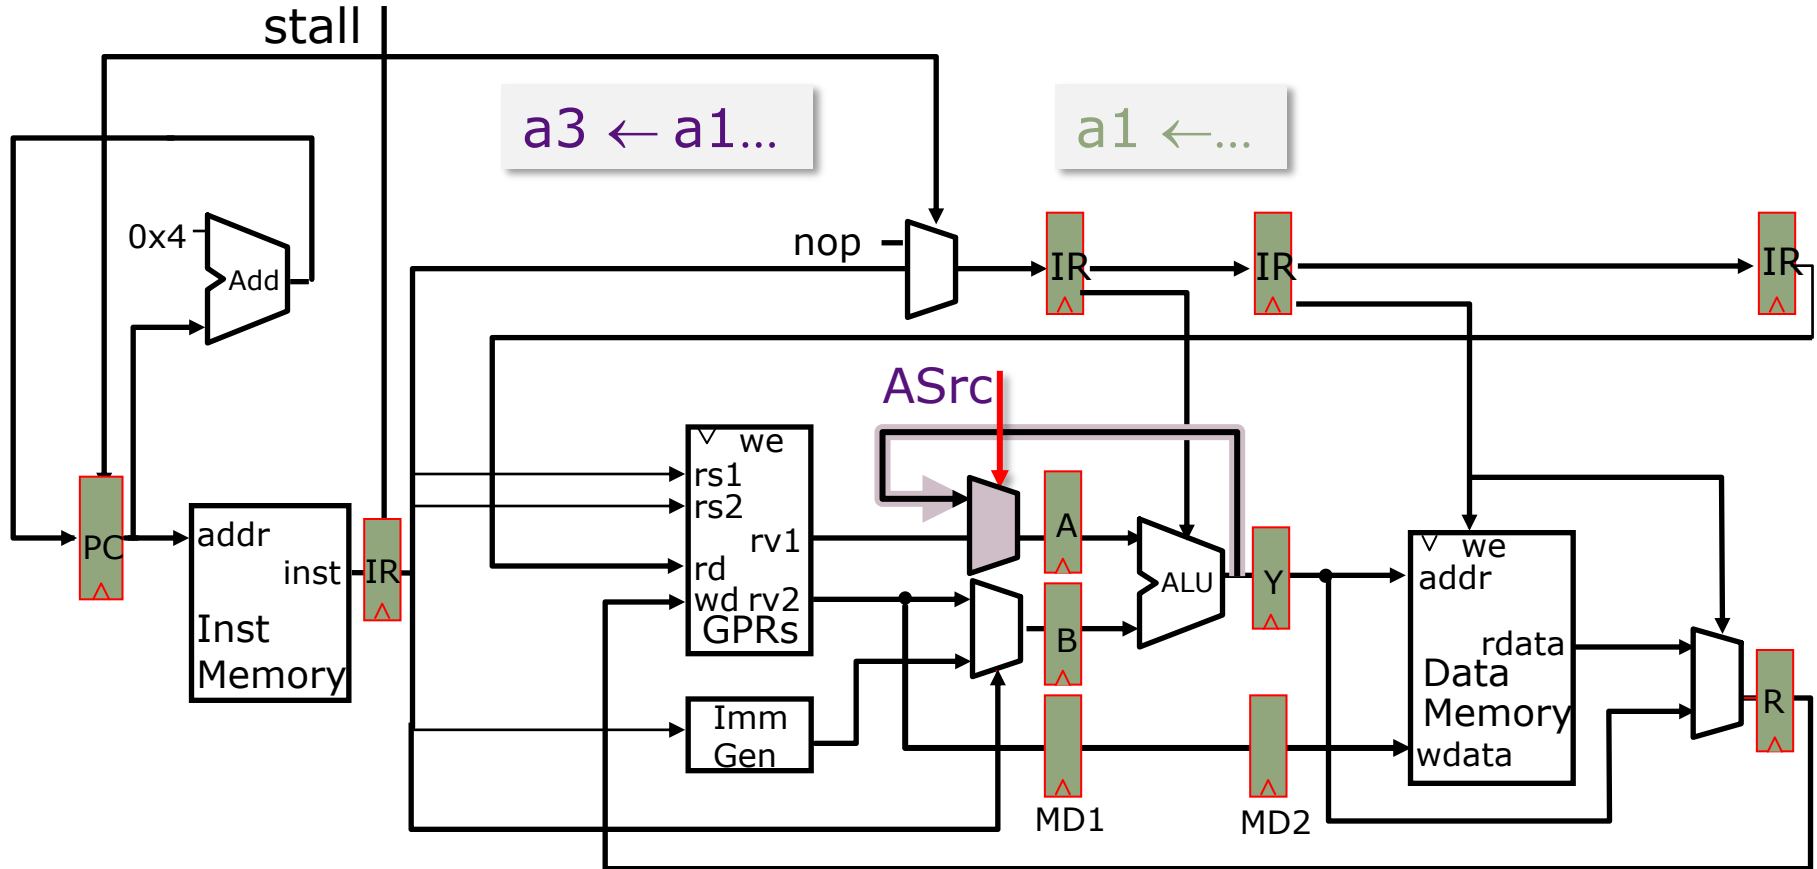
Each *stall* introduces a bubble $\Rightarrow CPI > 1$

When is data actually available?



A new datapath, i.e., a *bypass*, can get the data from the output of the ALU to its input

Adding a Bypass



When does this bypass help?

(I₁) a1 ← a0 + 10
(I₂) a3 ← a1 + 12 *yes*

a1 ← M[a0 + 10]
a3 ← a1 + 12 *no*

JAL ra 500
t3 ← ra + 12 *no*

The Bypass Signal

Deriving it from the Stall Signal

$$\text{stall} = \cancel{((rs1_D == rd_E) \cdot we_E)} + (rs1_D == rd_M) \cdot we_M + (rs1_D == rd_W) \cdot we_W \cdot re1_D \\ + ((rs2_D == rd_E) \cdot we_E + (rs2_D == rd_M) \cdot we_M + (rs2_D == rd_W) \cdot we_W) \cdot re2_D$$

we = Case opcode
ALU, ALUi, LW, JAL, JALR
 $\Rightarrow (rd \neq 0)$
... $\Rightarrow \text{off}$

$$\text{ASrc} = (rs1_D == rd_E) \cdot we_E \cdot re1_D$$

Is this correct?

How might we address this?

Bypass and Stall Signals

Split we_E into two components: we-bypass, we-stall

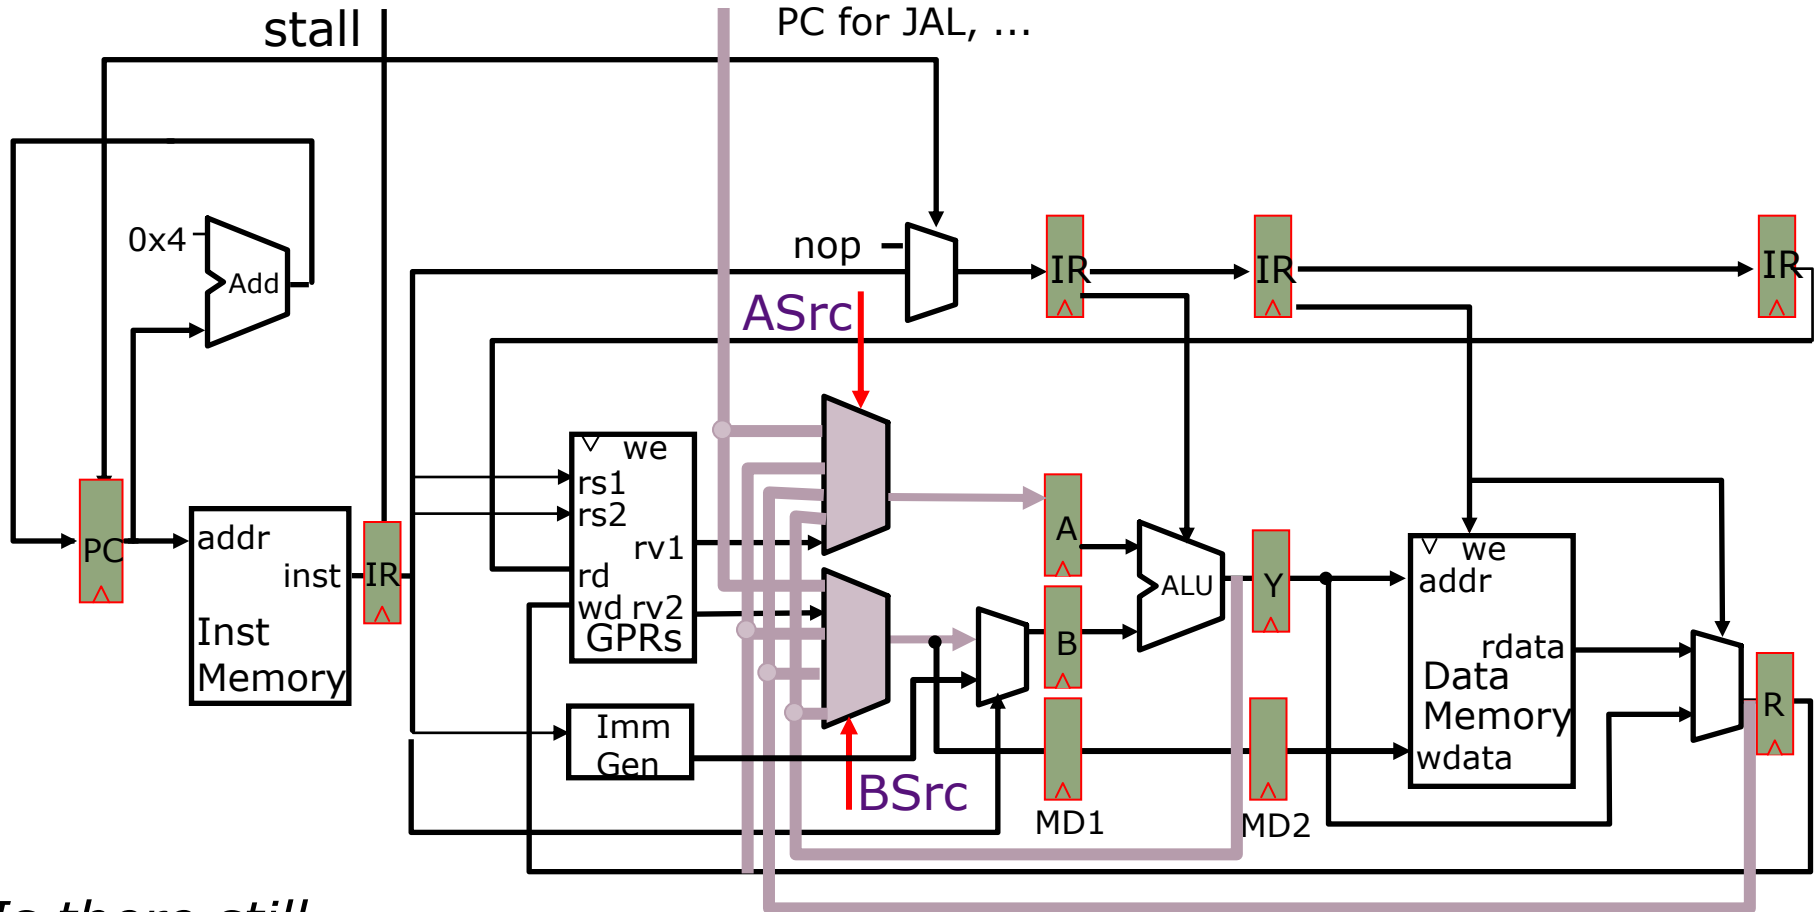
$we_bypass_E = \text{Case opcode}_E$
ALU, ALUi $\Rightarrow (rd \neq 0)$
... $\Rightarrow \text{off}$

$we_stall_E = \text{Case opcode}_E$
LW, JAL, JALR $\Rightarrow (rd \neq 0)$
... $\Rightarrow \text{off}$

$ASrc = (rs1_D == rd_E) \cdot re1_D \cdot we_bypass_E$

$stall = ((rs1_D == rd_E) \cdot we_stall_E +$
 $(rs1_D == rd_M) \cdot we_M + (rs1_D == rd_W) \cdot we_W) \cdot re1_D$
 $+ ((rs2_D == rd_E) \cdot we_E + (rs2_D == rd_M) \cdot we_M + (rs2_D == rd_W) \cdot we_W) \cdot re2_D$

Fully Bypassed Datapath



*Is there still
a need for the
stall signal?*

$$\text{stall} = (\text{rs1}_D == \text{rd}_E) \cdot (\text{opcode}_E == \text{LW}_E) \cdot (\text{rd}_E \neq 0) \cdot \text{re1}_D \\ + (\text{rs2}_D == \text{rd}_E) \cdot (\text{opcode}_E == \text{LW}_E) \cdot (\text{rd}_E \neq 0) \cdot \text{re2}_D$$

Resolving Data Hazards (3)

Strategy 3:

Speculate on the dependence. Two cases:

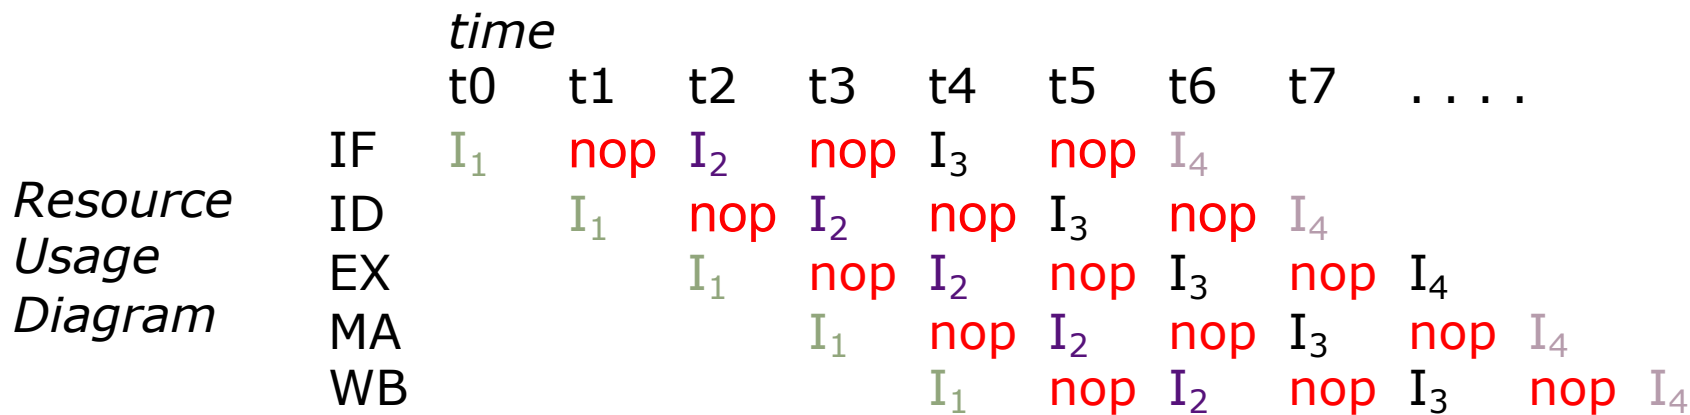

Guessed correctly → no special action required

Guessed incorrectly → kill and restart

Instruction to Instruction Dependence

- What do we need to calculate next PC?
 - For Jumps (JAL)
 - For Jump Register (JALR)
 - For Conditional Branches (e.g., BEQ)
 - For all others (e.g., arithmetic insts)
- In what stage do we know these?
 - PC →
 - Opcode, offset →
 - Register value →
 - Branch condition $((rs1) == (rs2))$ →

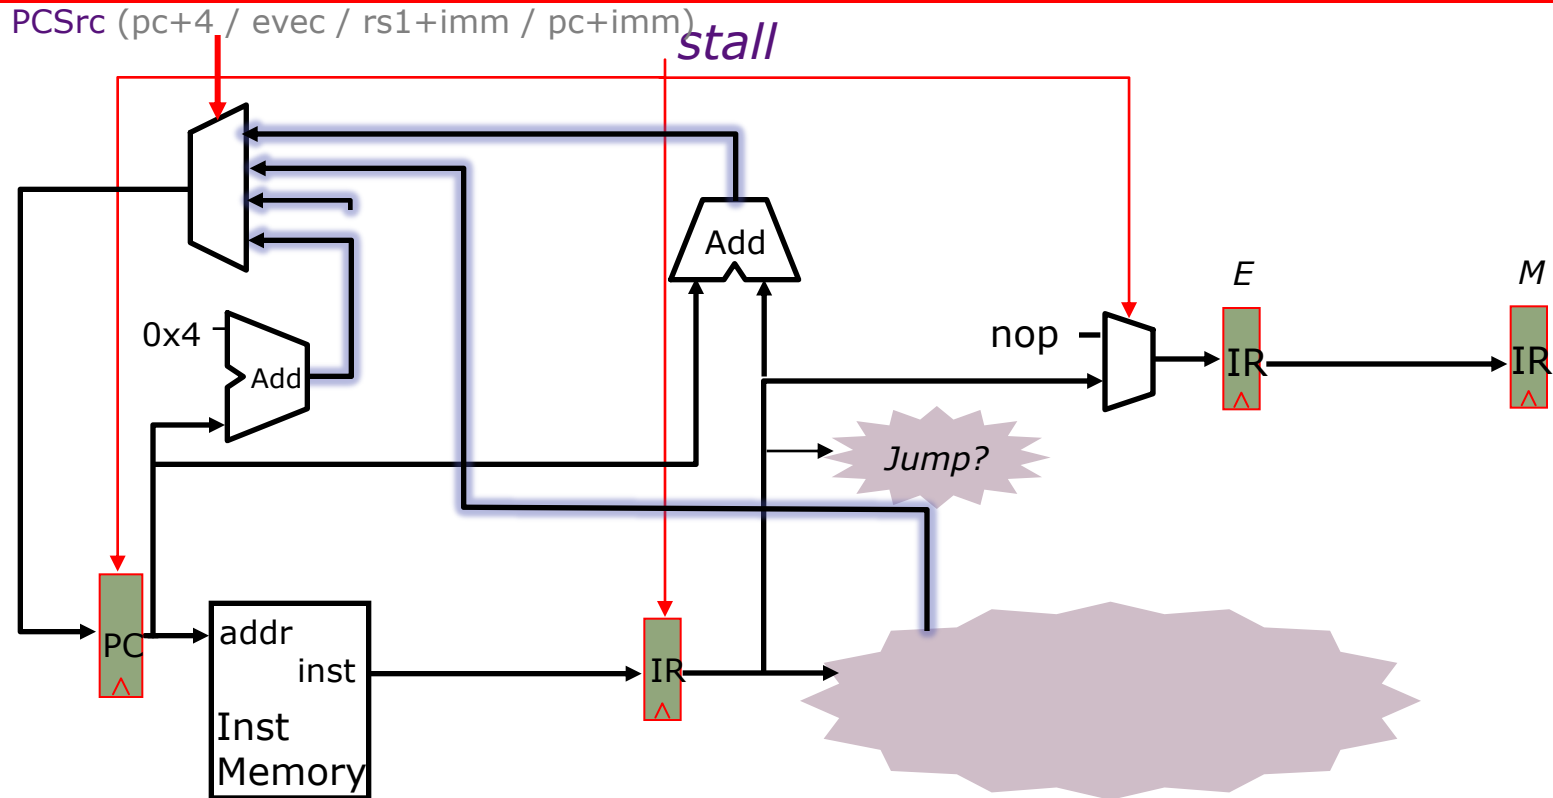
NextPC Calculation Bubbles



nop \Rightarrow *pipeline bubble*


What's a good guess for next PC?

Speculate NextPC is PC+4



Speculate NextPC is PC+4

PCSrc (pc+4 / evec / rs1+imm / pc+imm)



I ₁	096	ADD	
I ₂	100	JAL ra 200	
I ₃	104	ADD	<i>kill</i>
I ₄	304	ADD	

What happens on mis-speculation,
i.e., when next instruction is not PC+4?

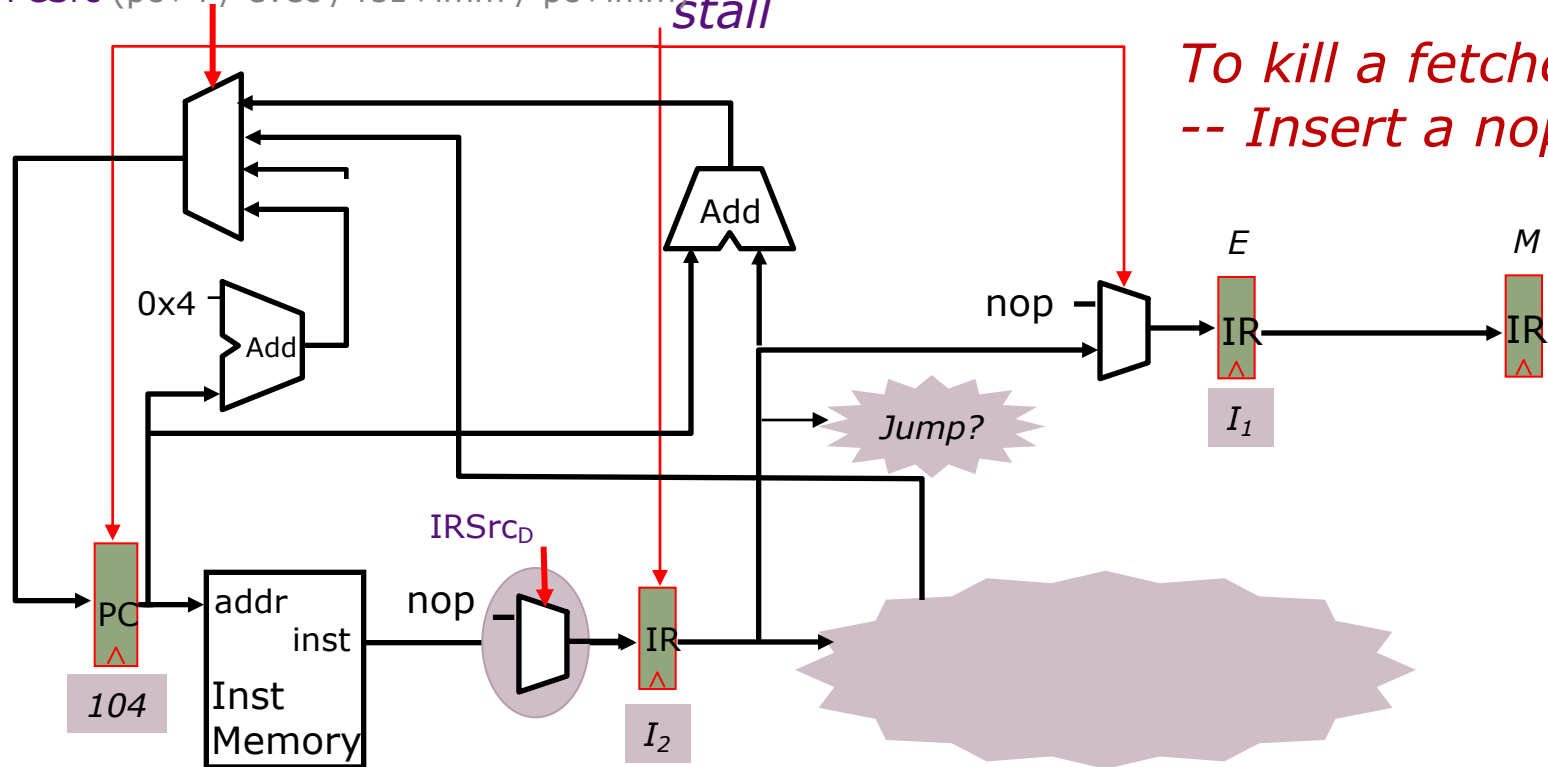
How?

Pipelining Jumps

PCSrc (pc+4 / evec / rs1+imm / pc+imm)

stall

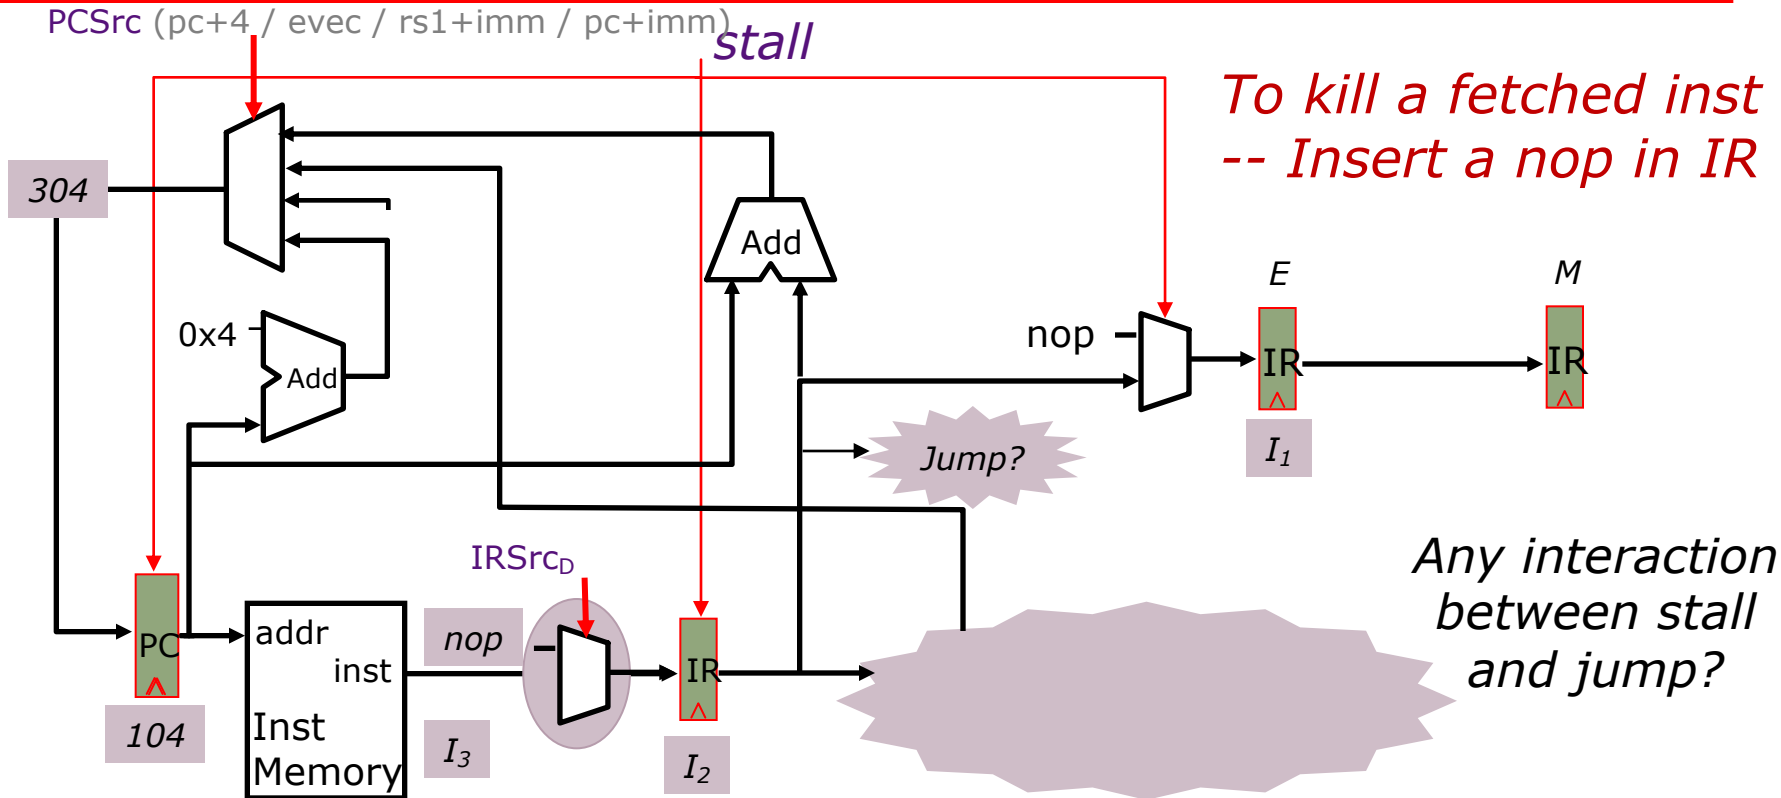
*To kill a fetched inst
-- Insert a nop in IR*



I_1	096	ADD	
I_2	100	JAL ra 200	
I_3	104	ADD	<i>kill</i>
I_4	304	ADD	

IRSrc_D = Case opcode_D
 JAL, JALR ⇒ nop
 ... ⇒ IM

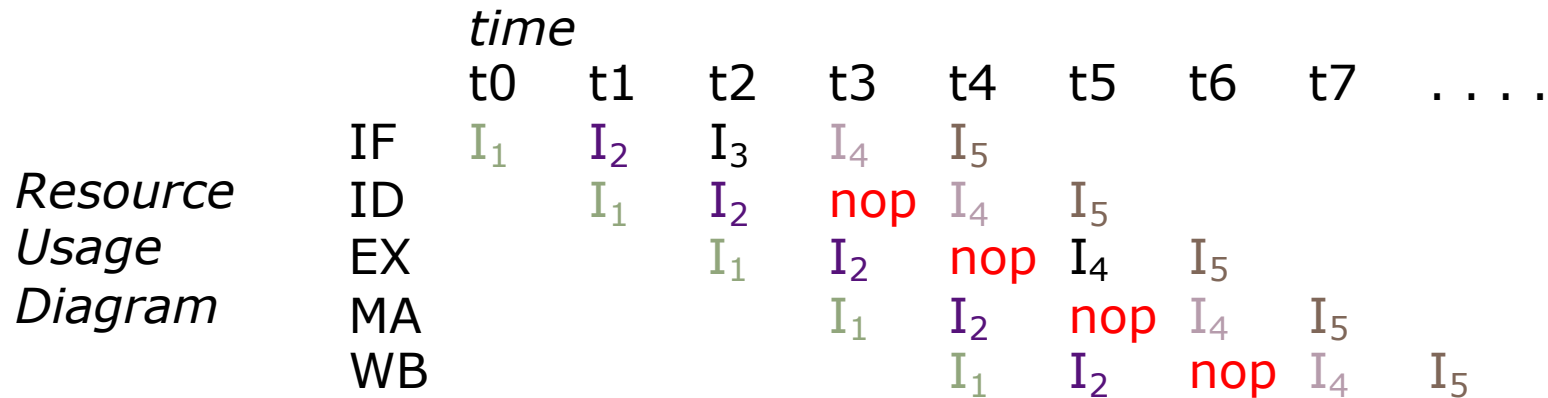
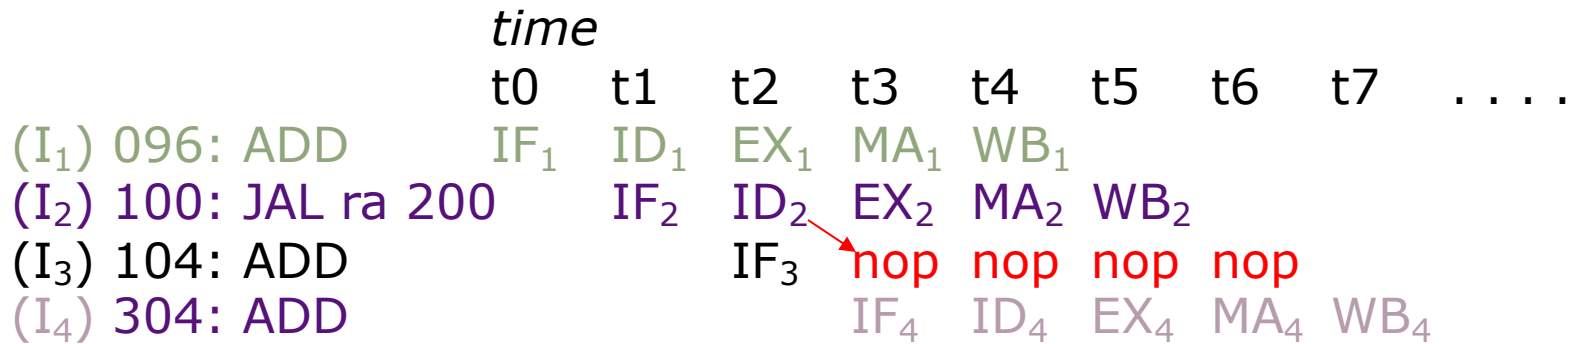
Pipelining Jumps



I_1	096	ADD	
I_2	100	JAL ra 200	
I_3	104	ADD	<i>kill</i>
I_4	304	ADD	

$IRSrc_D = \text{Case opcode}_D$
 JAL, JALR \Rightarrow nop
 ... \Rightarrow IM

Jump Pipeline Diagrams

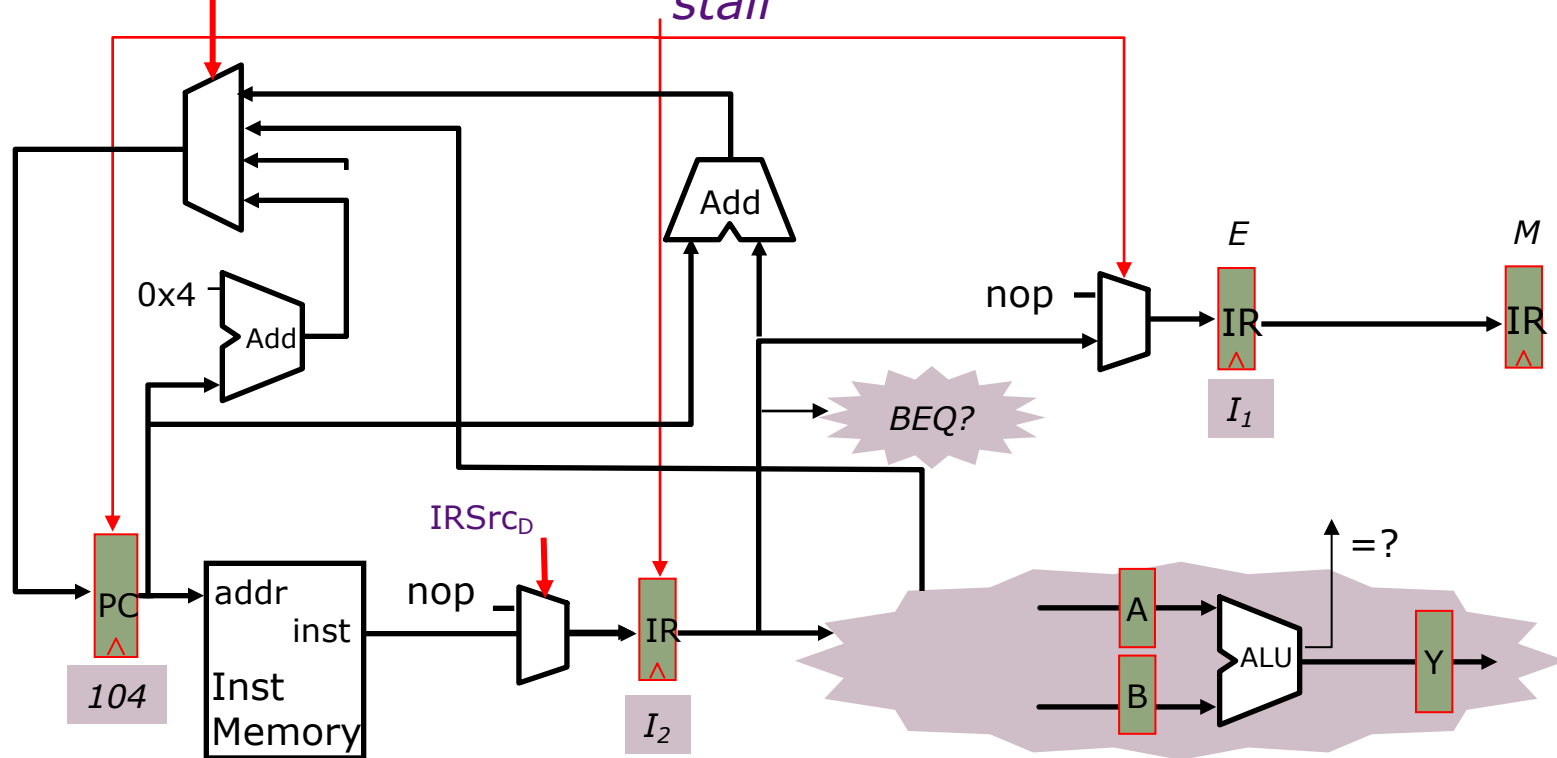


nop ⇒ *pipeline bubble*

Pipelining Conditional Branches

PCSrc (pc+4 / evec / rs1+imm / pc+imm)

stall



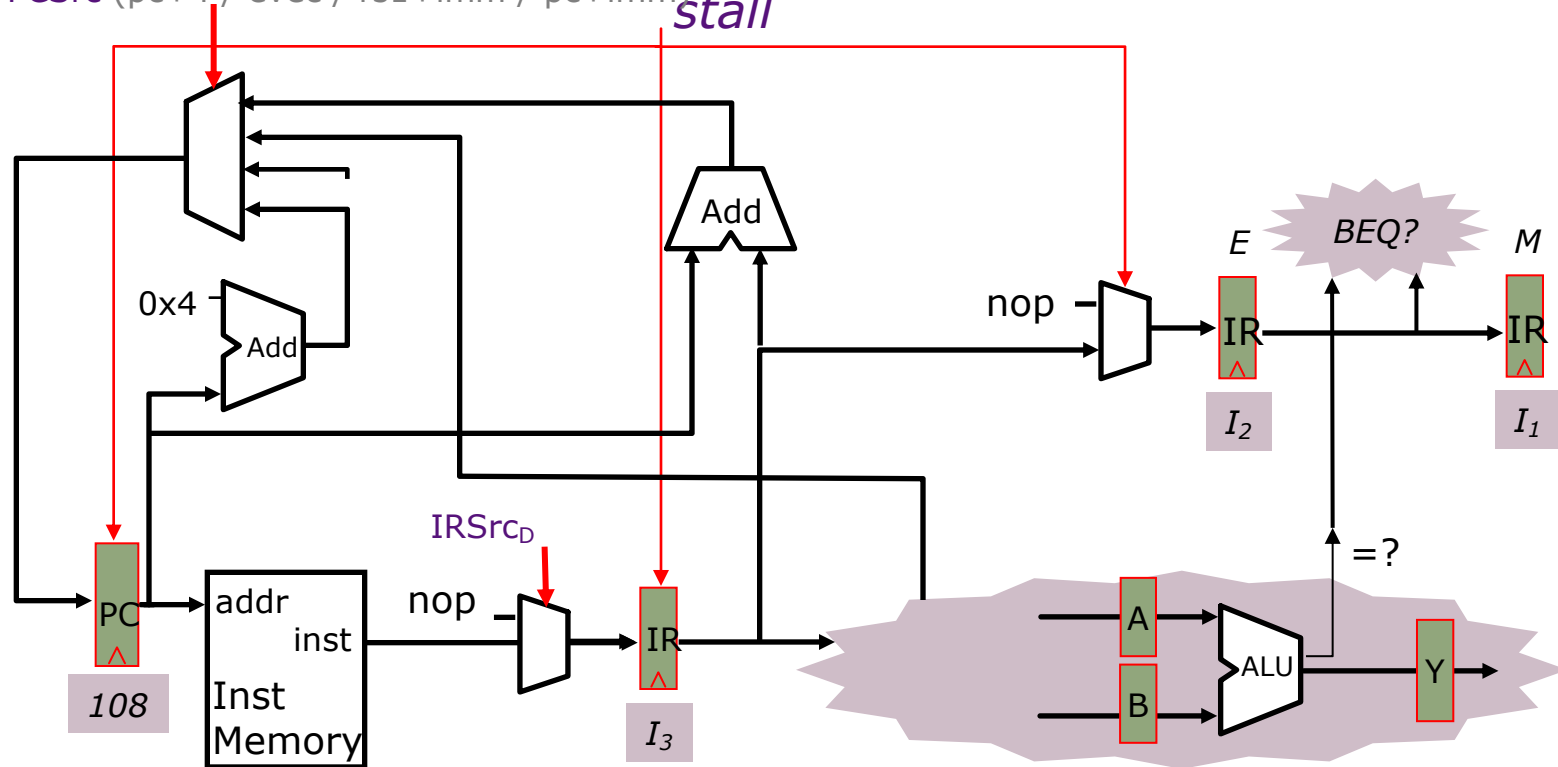
Branch condition is not known until the execute stage. *What action should be taken in the decode stage?*

I ₁	096	ADD
I ₂	100	BEQ a1 a2 200
I ₃	104	ADD
I ₄	304	ADD

Pipelining Conditional Branches

PCSrc (pc+4 / evec / rs1+imm / pc+imm)

stall



If the branch is taken

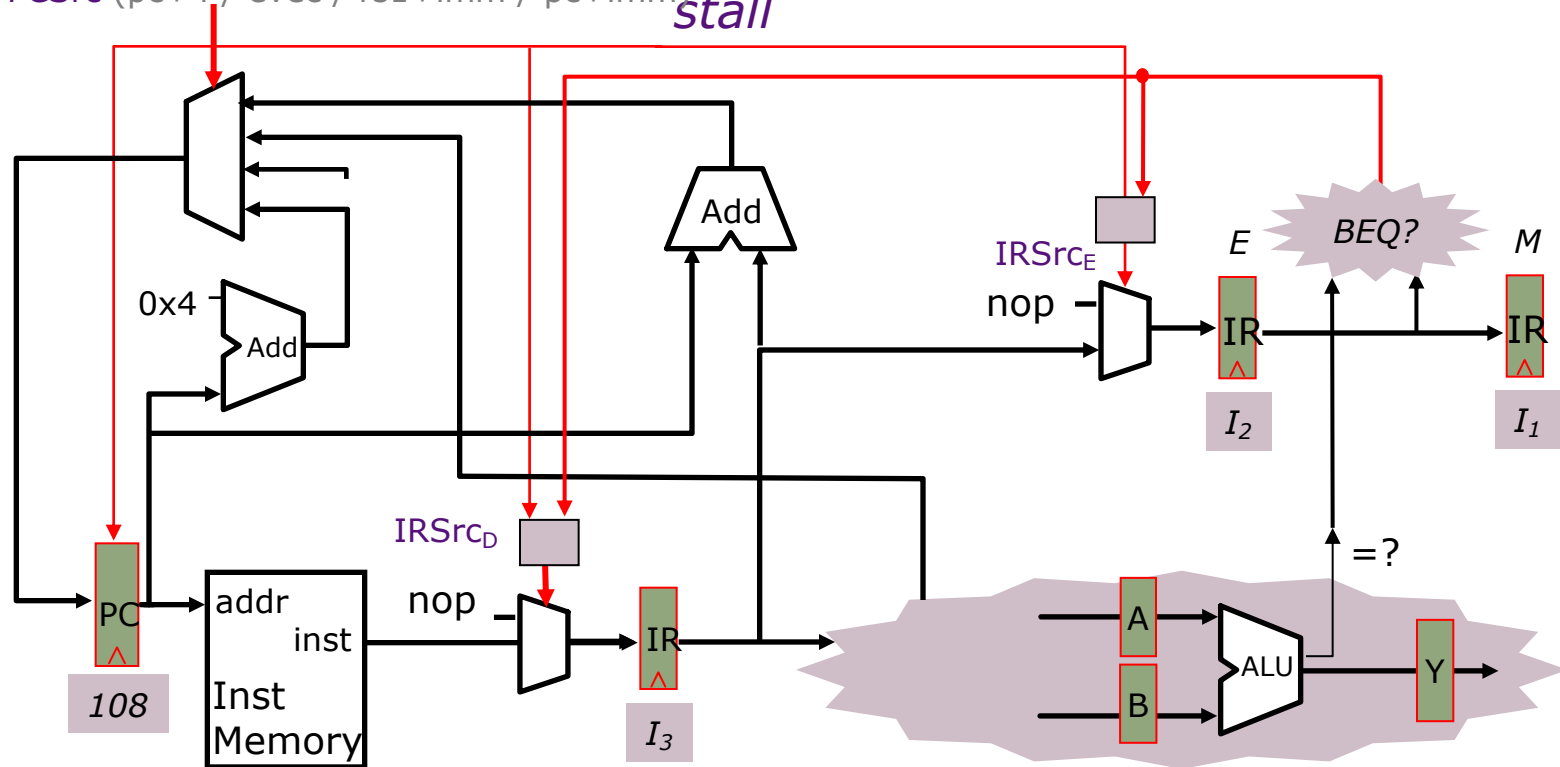
- kill the two following instructions
- the instruction at the decode stage is not valid

I ₁	096	ADD
I ₂	100	BEQ a1 a2 200
I ₃	104	ADD
I ₄	304	ADD

Pipelining Conditional Branches

PCSrc (pc+4 / evec / rs1+imm / pc+imm)

stall



I_1	096	ADD
I_2	100	BEQ a1 a2 200
I_3	104	ADD
I_4	304	ADD

Don't stall if the branch is taken.
Why?

Control Equations for PC and IR Muxes

$IRSrc_D = \text{Case opcode}_E$
Taken branch $\Rightarrow \text{nop}$
Case opcode_D
JAL, JALR $\Rightarrow \text{nop}$
... $\Rightarrow IM$

Give priority to the older instruction, i.e., execute stage instruction over decode stage instruction

$IRSrc_E = \text{Case opcode}_E$
Taken branch $\Rightarrow \text{nop}$
... $\Rightarrow \text{stall} \cdot \text{nop} + !\text{stall} \cdot IR_D$

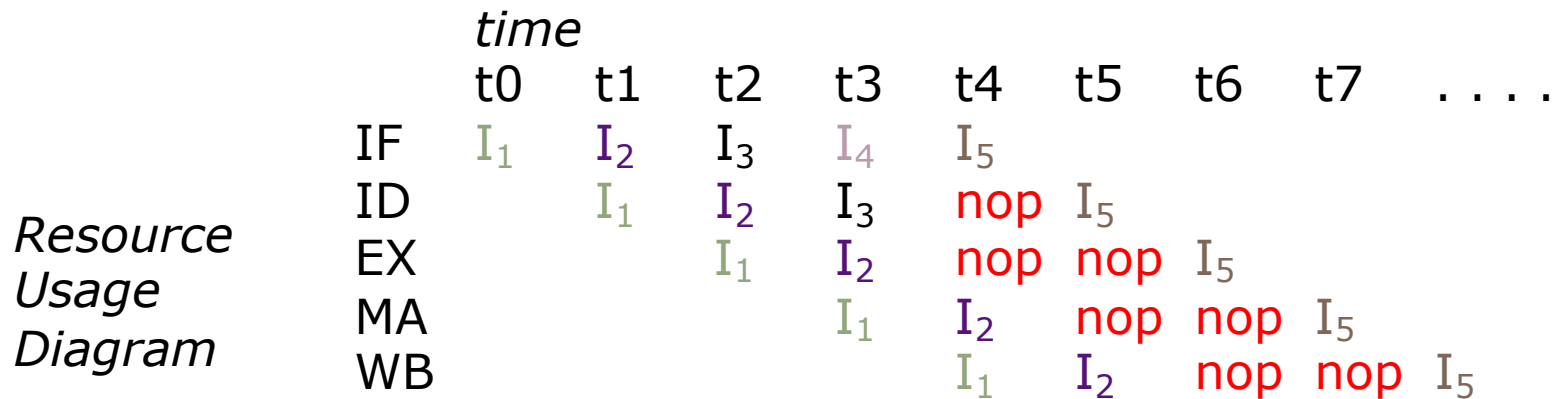
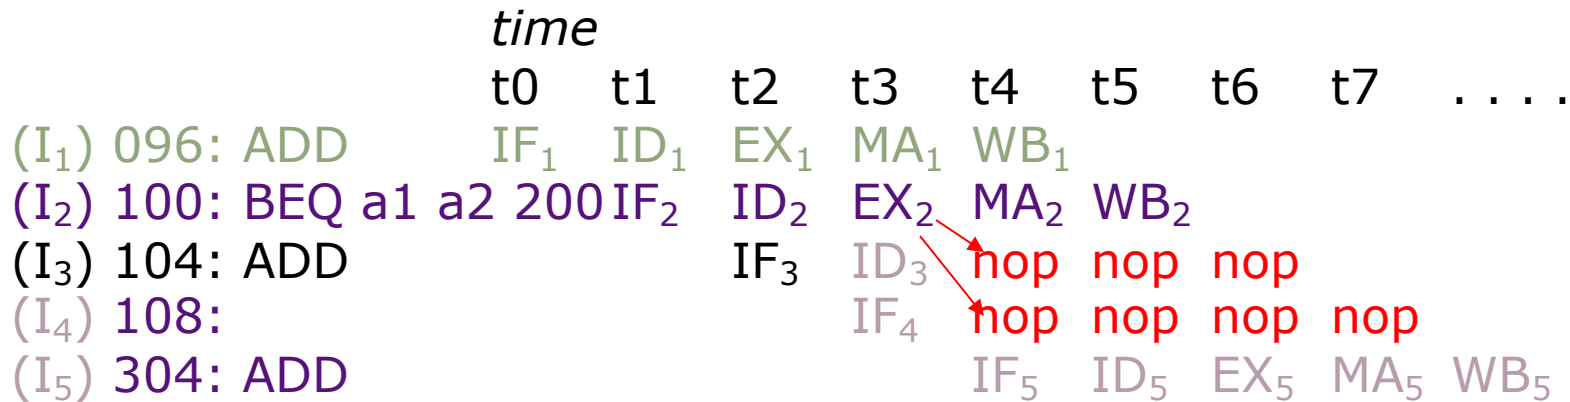
$PCSrc = \text{Case opcode}_E$
Taken branch $\Rightarrow pc + \text{imm (from E)}$
Case opcode_D
JAL $\Rightarrow pc + \text{imm (from D)}$
JALR $\Rightarrow rs1 + \text{imm (from D)}$
... $\Rightarrow pc+4$

pc+4 is a speculative guess

$\text{nop} \Rightarrow \text{Kill}$
 $pc+\text{imm} / rs1+\text{imm} \Rightarrow \text{Restart}$
 $pc+4 \Rightarrow \text{Speculate}$

Branch Pipeline Diagrams

(resolved in execute stage)

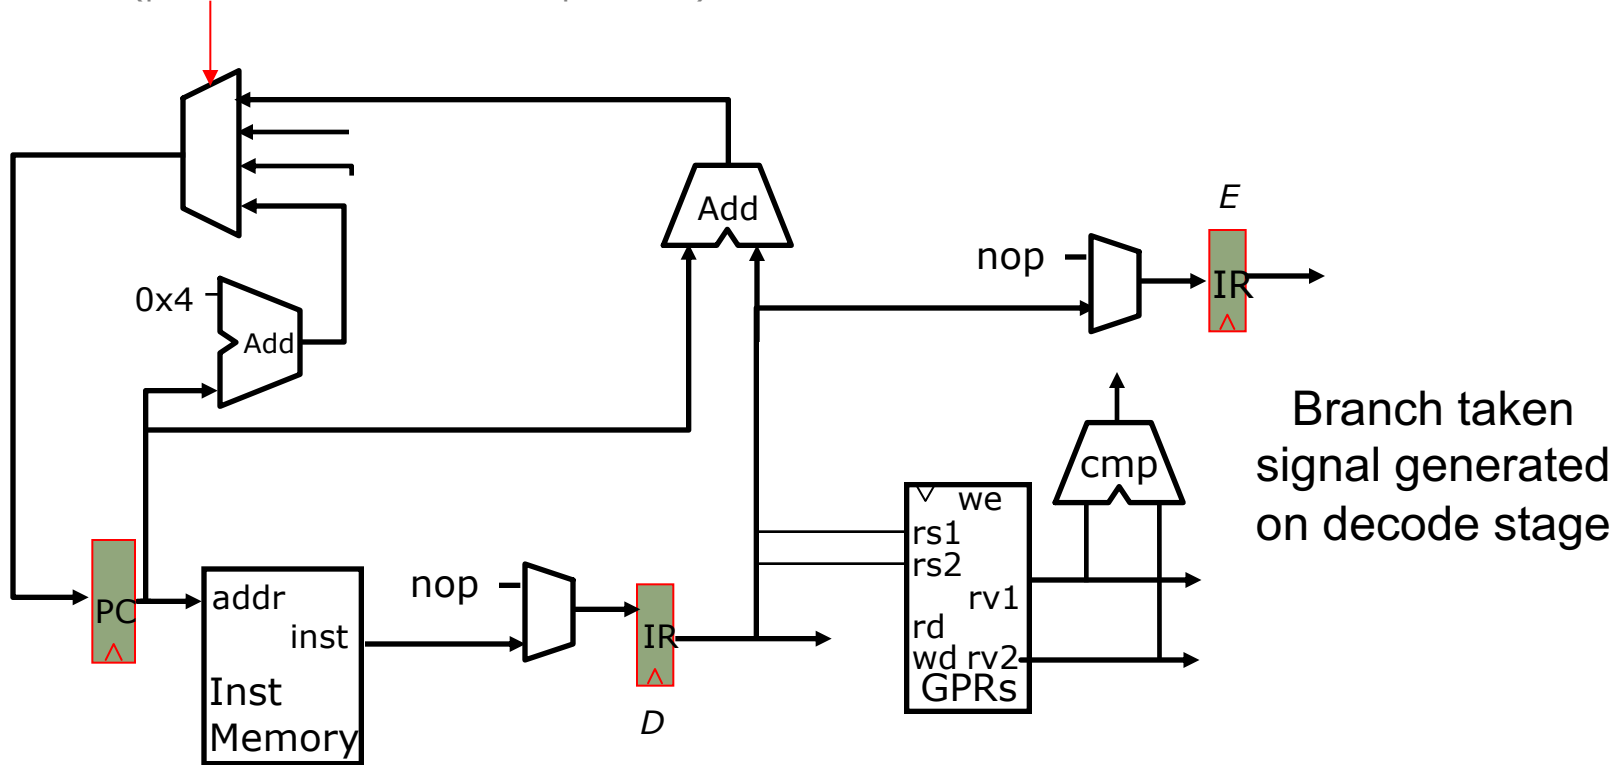


nop ⇒ *pipeline bubble*

Reducing Branch Penalty (resolve in decode stage)

- One pipeline bubble can be removed if an extra comparator is used in the Decode stage

PCSrc (pc+4 / evec / rs1+imm / pc+imm)



Pipeline diagram now same as for jumps

Branch Delay Slots

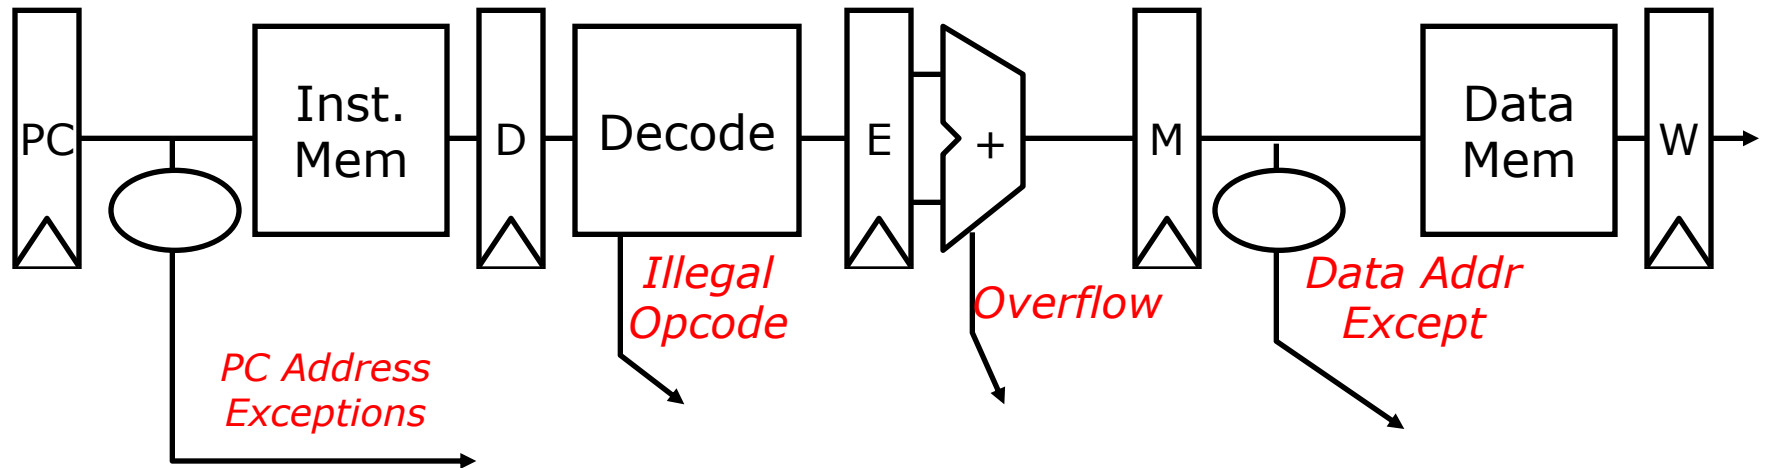
(expose control hazard to software)

- Change the ISA semantics so that the instruction that follows a jump or branch is always executed
 - gives compiler the flexibility to put in a useful instruction where normally a pipeline bubble would have resulted.

I ₁	096	ADD	
I ₂	100	BEQ a1 a2 200	<i>Delay slot instruction</i>
I ₃	104	ADD	← <i>executed regardless of</i>
I ₄	304	ADD	<i>branch outcome</i>

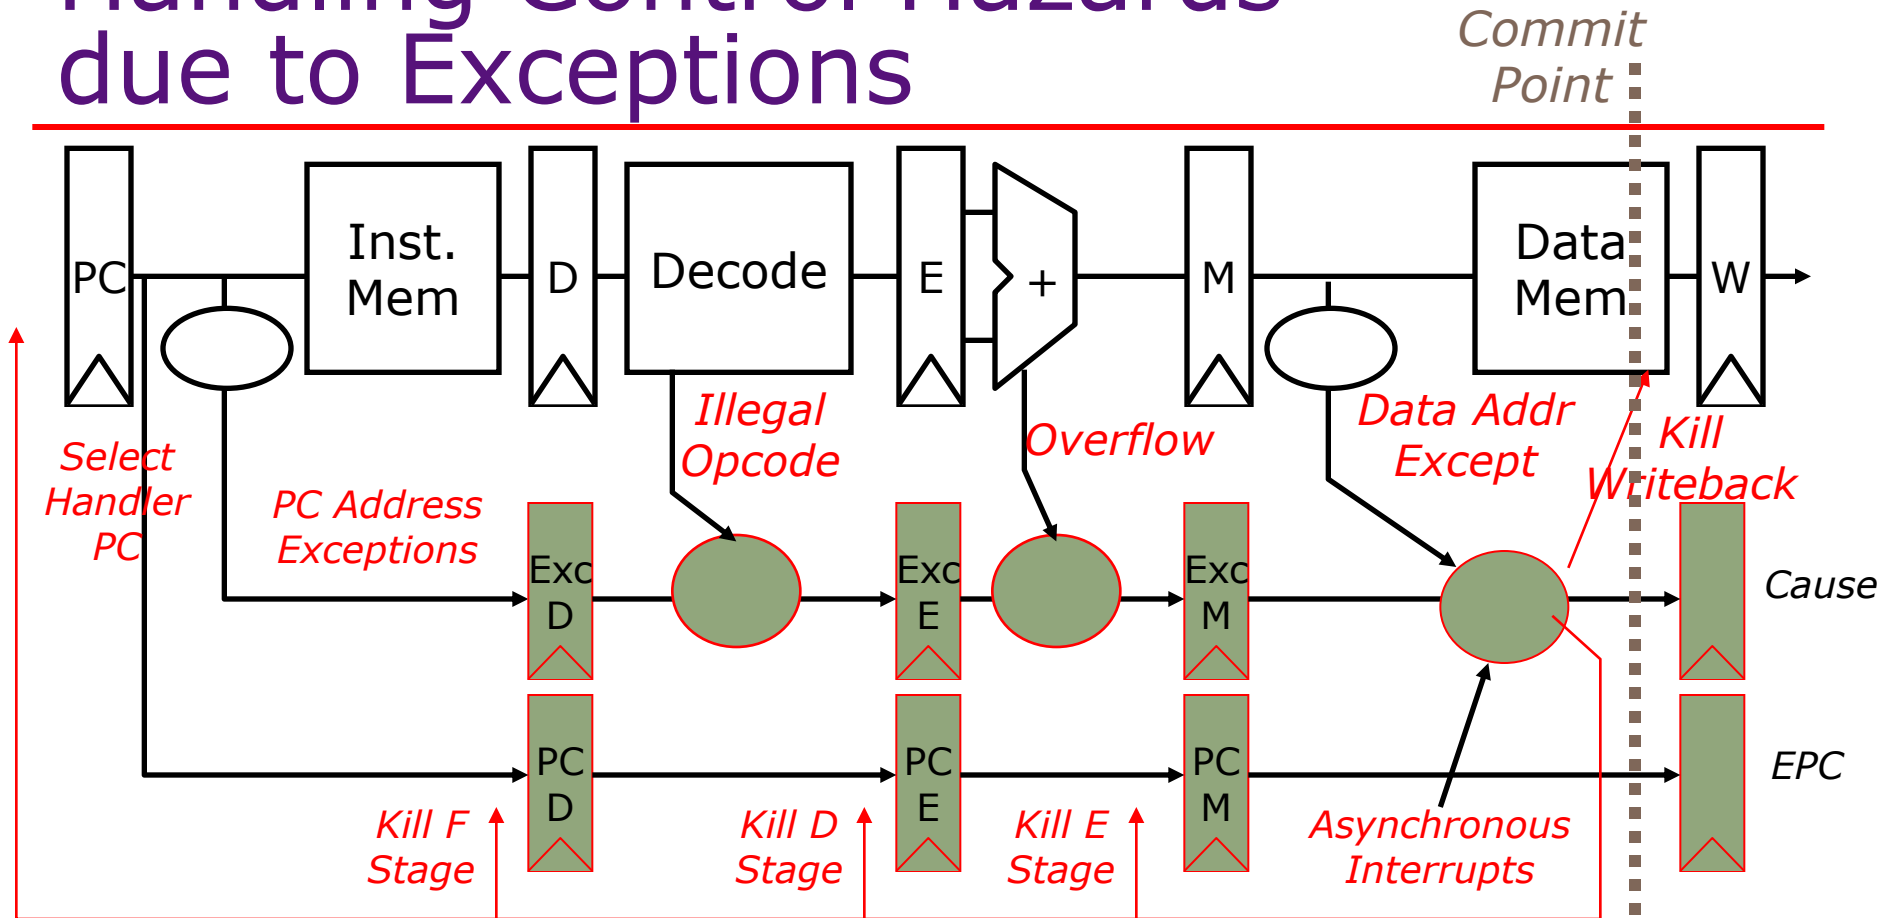
- Other techniques include branch prediction, which can dramatically reduce the branch penalty... *to come later*

Handling Control Hazards due to Exceptions



- Instructions may suffer exceptions in different pipeline stages
- Must prioritize exceptions from earlier instructions

Handling Control Hazards due to Exceptions



- Typical strategy: Record exceptions, process the first one to reach commit point (i.e., the point where architectural state is modified)
 - *Pros/cons vs handling exceptions eagerly, like branches?*

Why an instruction may not be dispatched every cycle ($CPI > 1$)

- Full bypassing may be too expensive to implement
 - Typically, all frequently used paths are provided
 - Some infrequently used bypass paths may increase cycle time and counteract the benefit of reducing CPI
- Loads have two-cycle latency
 - Instruction after load cannot use load result
 - MIPS-I ISA defined *load delay slots*, a software-visible pipeline hazard (compiler schedules independent instruction or inserts NOP to avoid hazard). Removed in MIPS-II.
- Conditional branches, jumps, and exceptions may cause bubbles
 - Kill instruction(s) following branch (if no delay slots, as in MIPS)

Machines with software-visible delay slots may execute significant number of NOP instructions inserted by the compiler

Next lecture:
Superscalar & Scoreboarded
Pipelines