Directory-Based Cache Coherence

Daniel Sanchez
Computer Science and Artificial Intelligence Lab
M.I.T.

Maintaining Cache Coherence

It is sufficient to have hardware such that

• Only one processor at a time has write permission for a location

• No processor can load a stale copy of the location after a write

$\Rightarrow$ A correct approach could be:

write request:
The address is *invalidated* in all other caches *before* the write is performed

read request:
If a dirty copy is found in some cache, a write-back is performed before the memory is read

Directory-Based Coherence
[Censier and Feautrier, 1978]

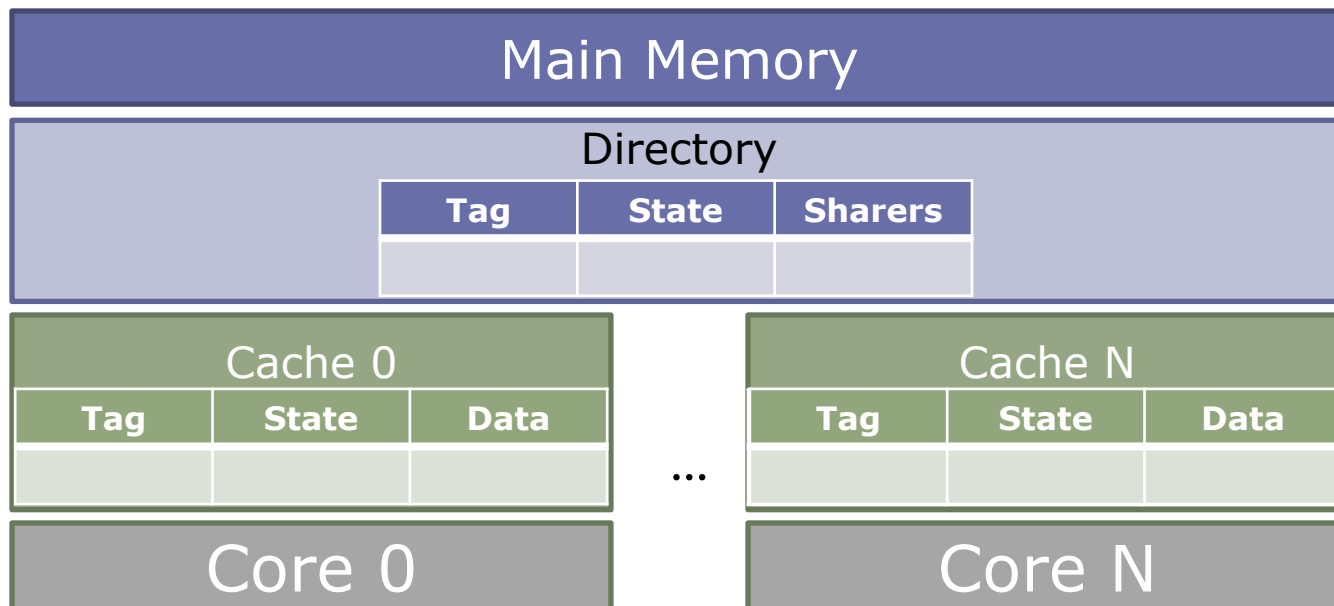Snoopy Protocols

Directory Protocols

- Snoopy schemes broadcast requests over memory bus

- Difficult to scale to large numbers of processors

- Requires additional bandwidth to cache tags for snoop requests

- Directory schemes send messages to only those caches that might have the line

- Can scale to large numbers of processors

- Requires extra directory storage to track possible sharers

An MSI Directory Protocol

Main Memory		

Directory

Tag	State	Sharers

Cache 0		

Tag	State	Data

...

Cache N		

Tag	State	Data

Core 0

Core N

- Cache states: Modified (M) / Shared (S) / Invalid (I)
- Directory states:
  - Uncached (Un): No sharers
  - Shared (Sh): One or more sharers with read permission (S)
  - Exclusive (Ex): A single sharer with read & write permissions (M)
- Transient states not drawn for clarity; for now, assume no racing requests

MSI Protocol: Caches (1/3)

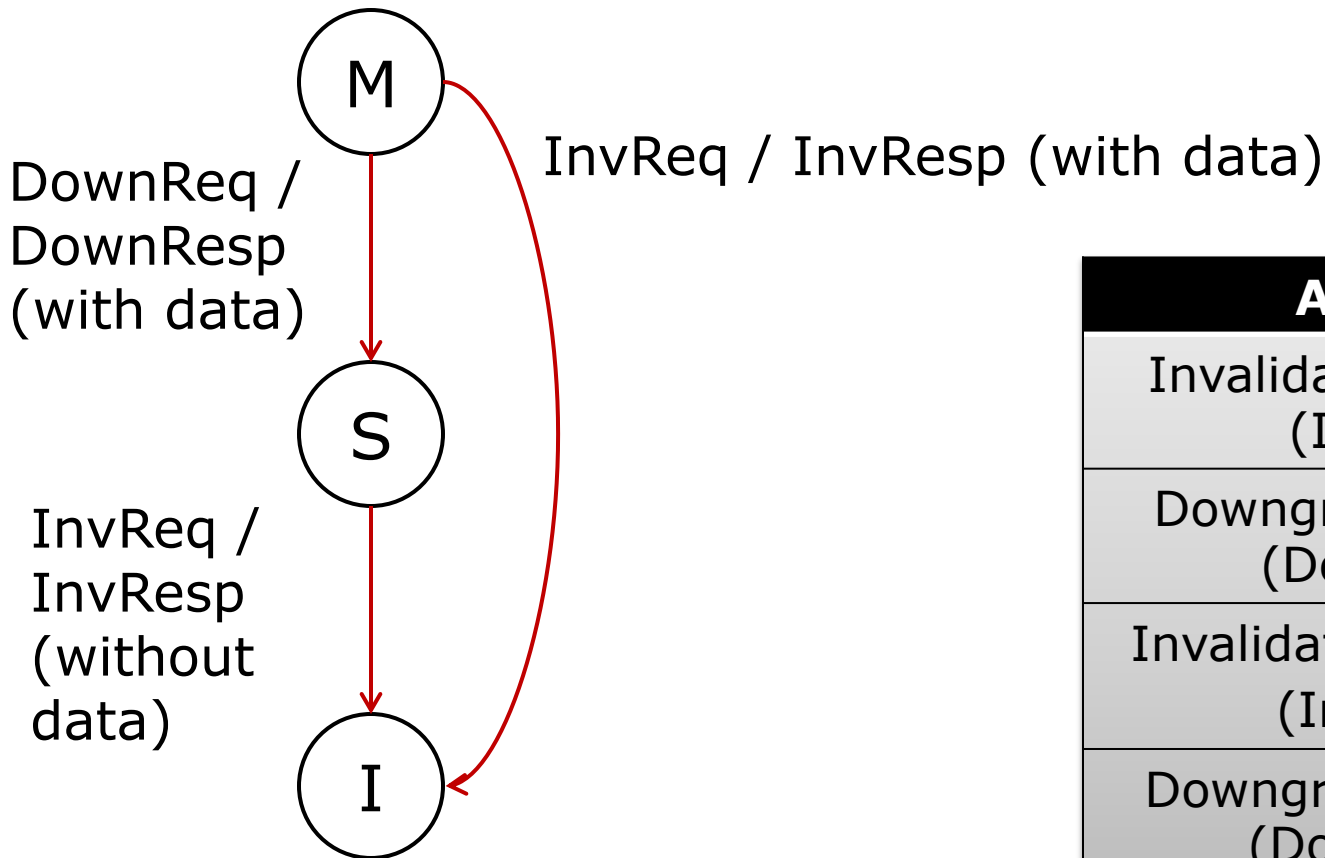Transitions initiated by processor accesses:

PrRd / --
PrWr / --

M

PrWr / ExReq

PrWr / ExReq

S

PrRd / --

PrRd / ShReq

I

Actions
Processor Read (PrRd)
Processor Write (PrWr)
Shared Request (ShReq)
Exclusive Request (ExReq)

MSI Protocol: Caches (2/3)

Transitions initiated by directory requests:

Actions
Invalidation Request (InvReq)
Downgrade Request (DownReq)
Invalidation Response (InvResp)
Downgrade Response (DownResp)

State diagram:
- M → S: DownReq / DownResp (with data)
- M → I: InvReq / InvResp (with data)
- S → I: InvReq / InvResp (without data)

MSI Protocol: Caches (3/3)

Transitions initiated by evictions:

M

Eviction /
WbReq
(with data)

S

Eviction /
WbReq
(without data)

I

Actions
Writeback Request (WbReq)

MSI Protocol: Caches

$\longrightarrow$ Transitions initiated by processor accesses

$\longrightarrow$ Transitions initiated by directory requests

$\longrightarrow$ Transitions initiated by evictions

MSI Protocol: Directory (1/2)

Transitions initiated by data requests:

ExReq / Sharers = {P}; ExResp

ExReq / Inv(Sharers); Sharers={P}; ExResp

ShReq / Down(Sharer); Sharers = Sharer + {P}; ShResp

ExReq / Inv(Sharers – {P}); Sharers = {P}; ExResp

ShReq / Sharers = Sharers + {P}; ShResp

ShReq / Sharers = {P}; ShResp

(States: Ex, Sh, Un)
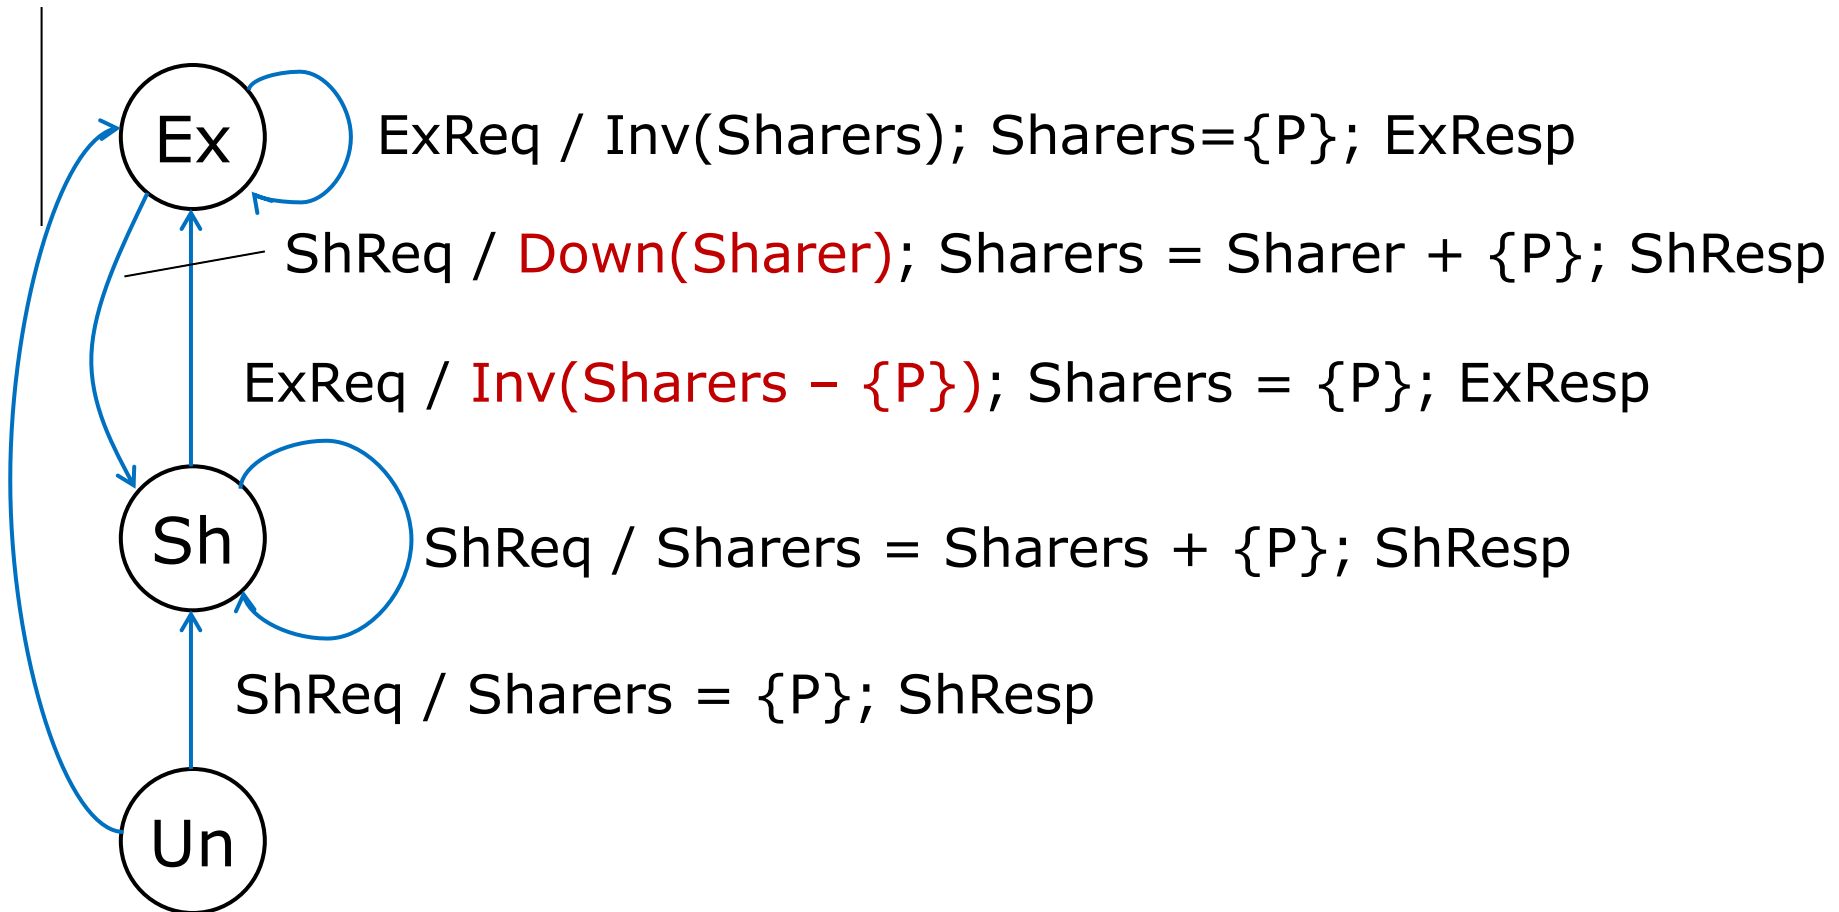
MSI Protocol: Directory (2/2)

Transitions initiated by writeback requests:

Ex

WbReq / Sharers = {}; WbResp

Sh

WbReq && |Sharers| > 1 /
Sharers = Sharers - {P}; WbResp

WbReq && |Sharers| == 1 /
Sharers = {}; WbResp

Un

MSI Directory Protocol Example

Main Memory

③ Mem[0xA] = 3

Directory

Tag	State	Sharers
0xA	Sh	{0}

② ShReq 0xA

④ ShResp 0xA, data=3

Cache 0

Tag	State	Data
0xA	S	3

Cache 1

Tag	State	Data

Cache 2

Tag	State	Data

Core 0

Core 1

Core 2

① LD 0xA

MSI Directory Protocol Example

Main Memory

③ Mem[0xA] = 3

Directory

Tag	State	Sharers
0xA	Sh	{0,2}

④ ShResp 0xA, data=3

② ShReq 0xA

Cache 0

Tag	State	Data
0xA	S	3

Cache 1

Tag	State	Data

Cache 2

Tag	State	Data
0xA	S	3

Core 0

Core 1

Core 2

① LD 0xA

MSI Directory Protocol Example

MSI Directory Protocol Example

Main Memory

③ Mem[0xA] = 5   **⑥** Mem[0xB] = 10

Directory

Tag	State	Sharers
0xB	Ex	{1}

② WbReq 0xA, data=5   **⑤** ExReq 0xB

④ WbResp 0xA   **⑦** ExResp 0xB, data=10

Cache 0

Tag	State	Data
0xA	I	3

Cache 1

Tag	State	Data
0xB	M	10

Cache 2

Tag	State	Data
0xA	I	3

Core 0   **Core 1**   **Core 2**

① ST 0xB
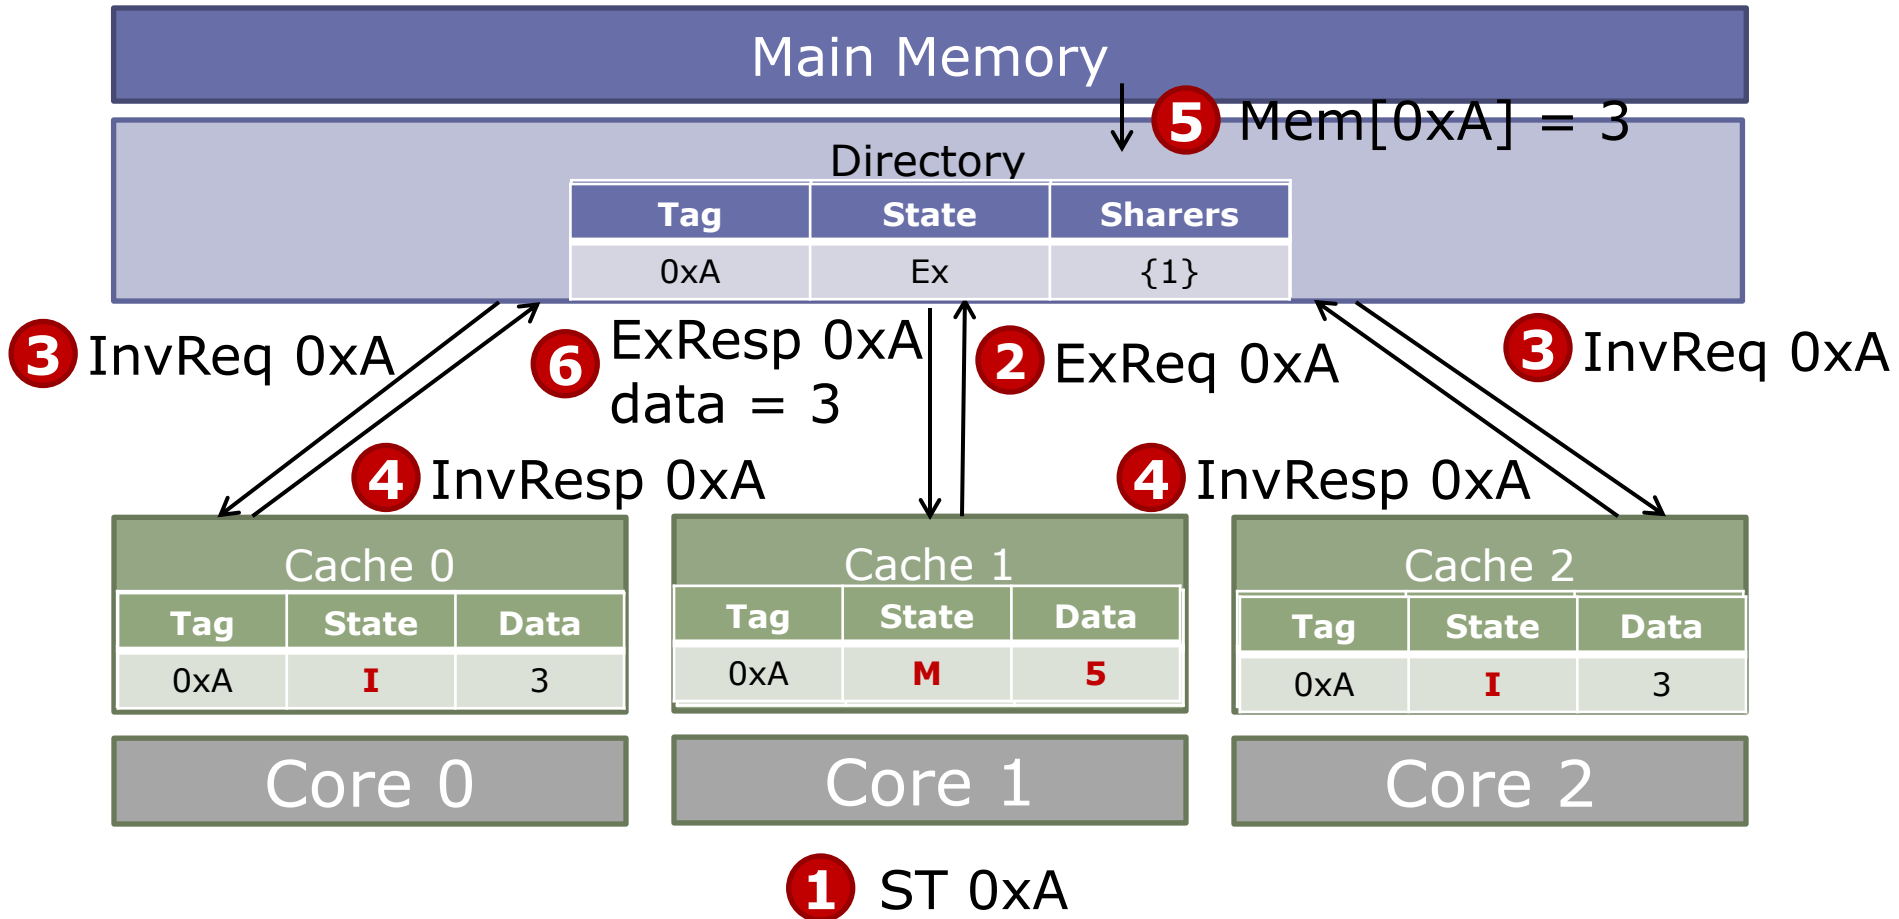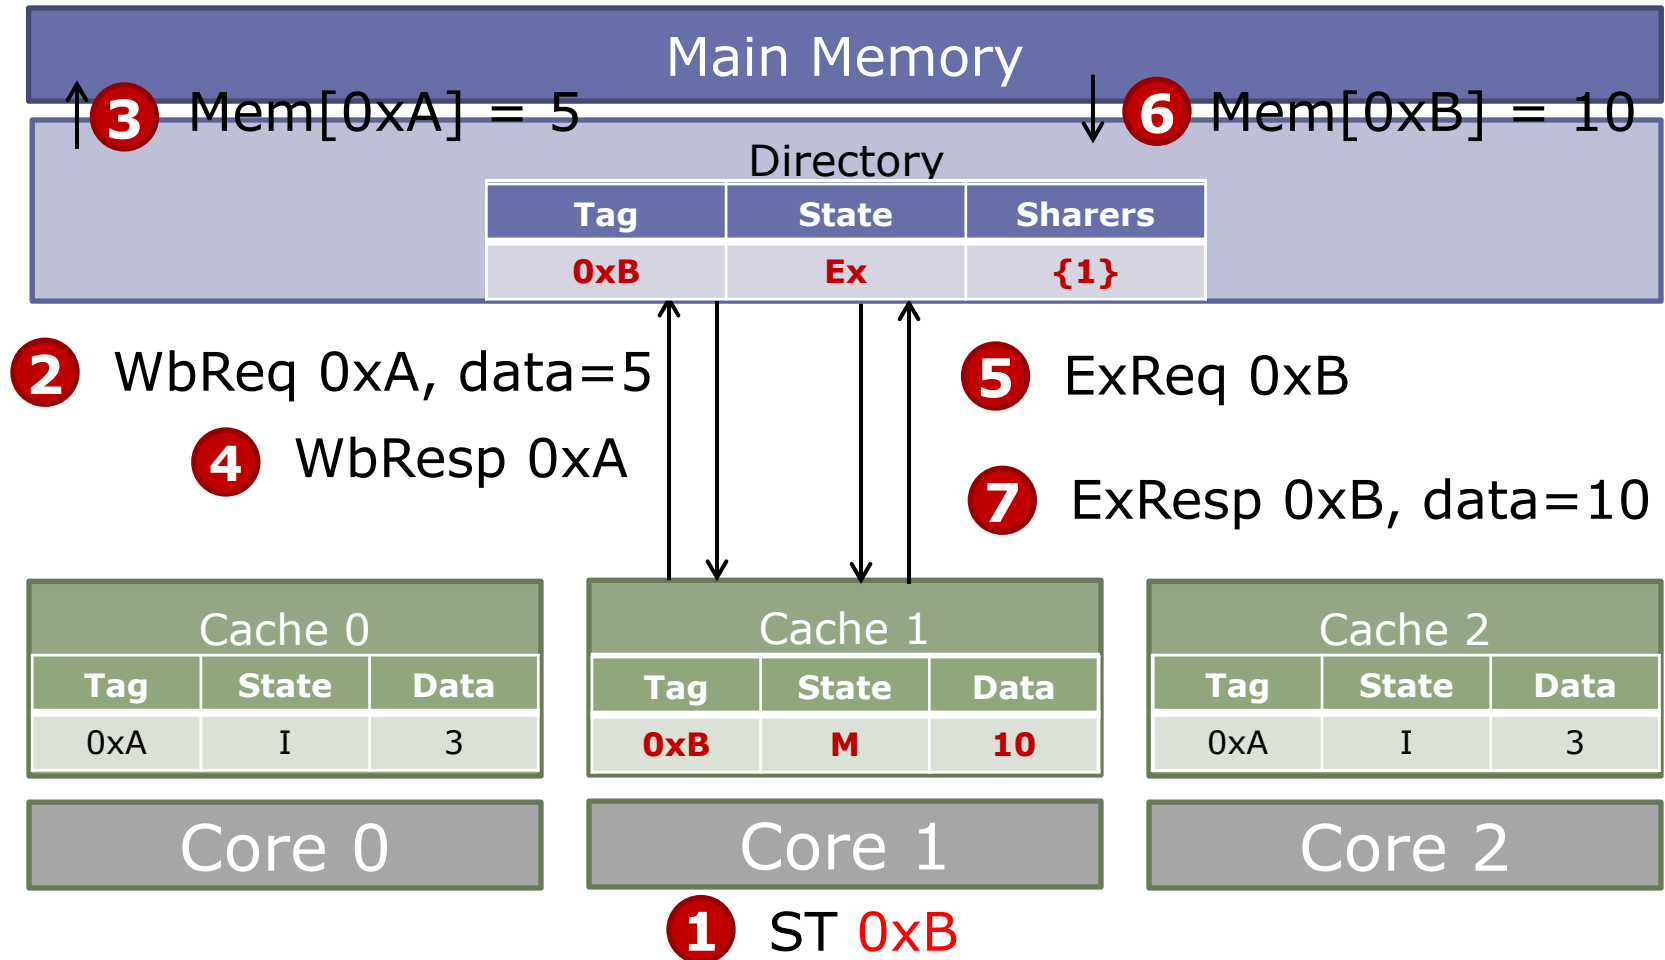
Why are 0xA's wb and 0xB's req serialized?
Possible solutions?

Miss Status Holding Register

MSHR – Holds load misses and writes outside of cache

MSHR entry

V	X	Addr	Data

- ## On eviction/writeback
  - No free MSHR entry: stall
  - Allocate new MSHR entry
  - When channel available send WBReq and data
  - Deallocate entry on WBResp
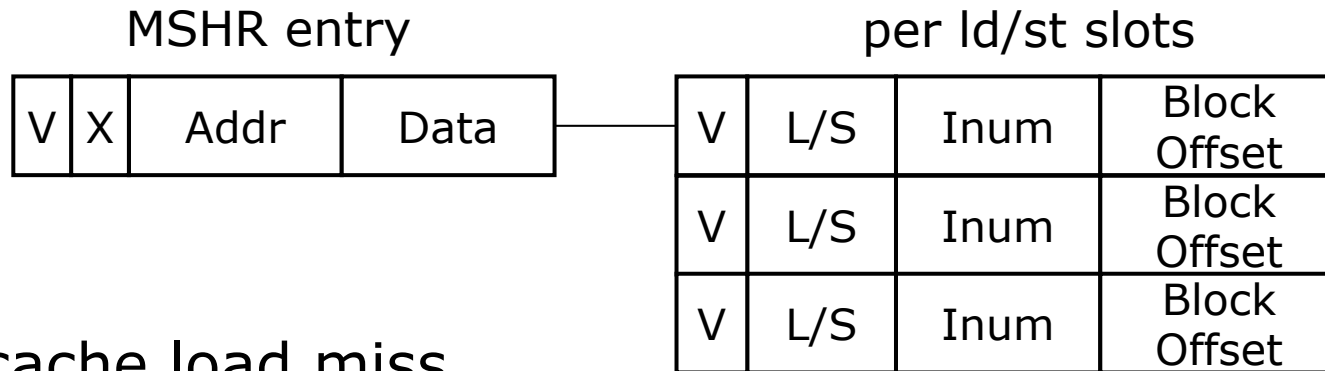
Miss Status Holding Register

MSHR – Holds load misses and writes outside of cache

<div align="center">

MSHR entry   per ld/st slots

V	X	Addr	Data		L/S	Inum	Block Offset

</div>

- ## On cache load miss
  - No free MSHR entry: stall
  - Allocate new MSHR entry
  - Send ShReq (or ExReq)
  - On *Resp forward data to CPU and cache
  - Deallocate MSHR

Miss Status Holding Register

MSHR – Holds load misses and writes outside of cache

MSHR entry				per ld/st slots			
V	X	Addr	Data	V	L/S	Inum	Block Offset
				V	L/S	Inum	Block Offset
				V	L/S	Inum	Block Offset

- ## On cache load miss
  - Look for matching address in MSHRs
    - If not found
      - If no free MSHR entry: stall
      - Allocate new MSHR entry and fill in
    - If found, just fill in per ld/st slot
  - Send ShReq (or ExReq)
  - On *Resp forward data to CPU and cache
 - Deallocate MSHR

Per ld/st slots allow servicing multiple requests with one entry

Directory Organization

- Requirement: Directory needs to keep track of all the cores that are sharing a cache block

- Challenge: For each block, the space needed to hold the list of sharers grows with number of possible sharers…

Flat, Memory-based Directories

- Dedicate a few bits of main memory to store the state and sharers of every line

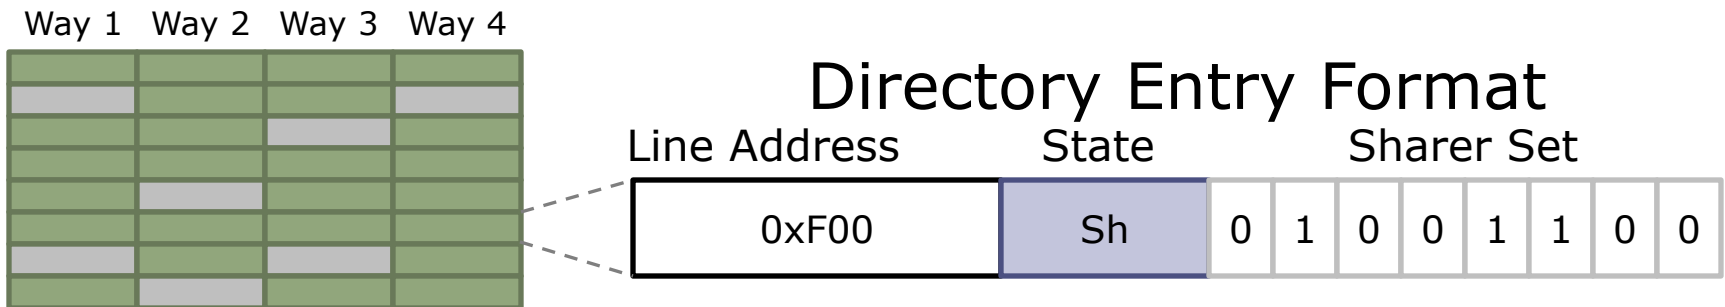- Encode sharers using a bit-vector

Main Memory

State	Sharer Set							
Sh	0	1	0	0	1	1	0	0

64 bytes   10 bits

✓ Simple
✗ Slow
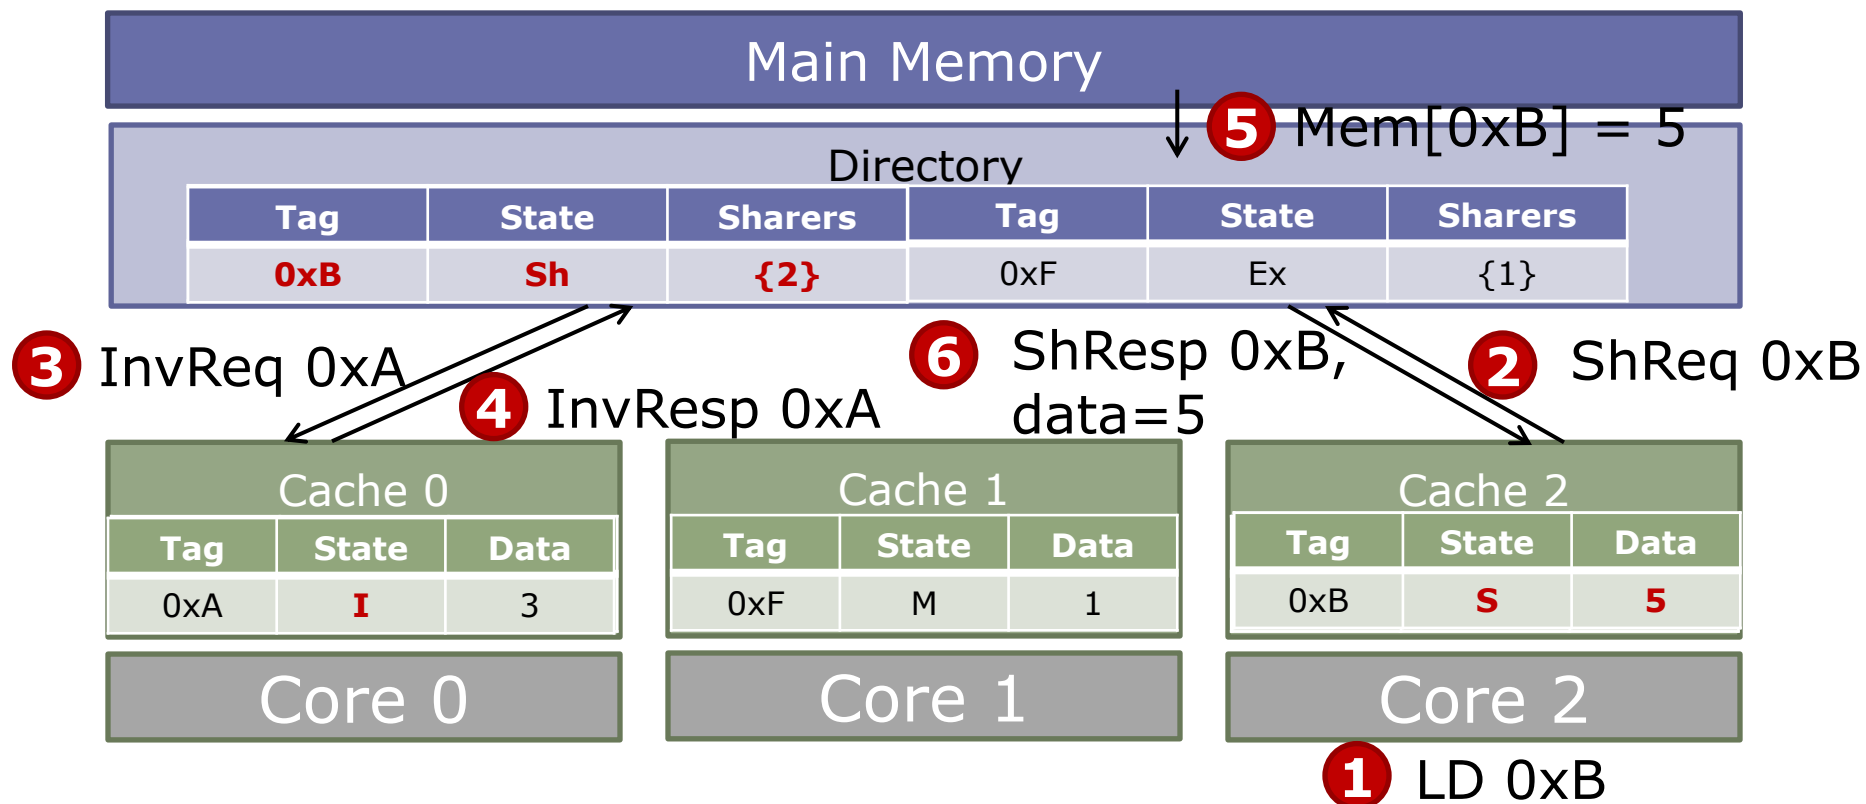✗ Very inefficient with many processors (~P bits/line)

Sparse Full-Map Directories

- Not every line in the system needs to be tracked – only those in private caches!

- Idea: Organize directory as a cache

Way 1 Way 2 Way 3 Way 4

Directory Entry Format

Line Address	State	Sharer Set							
0xF00	Sh	0	1	0	0	1	1	0	0

✓ Low latency, energy-efficient
✗ Bit-vectors grow with # cores → Area scales poorly
✗ Limited associativity → Directory-induced invalidations

Directory-Induced Invalidations

- To retain **inclusion**, must invalidate all sharers of an entry before reusing it for another address

- Example: 2-way set-associative sparse directory

Main Memory					

Directory

Tag	State	Sharers	Tag	State	Sharers
0xB	**Sh**	**{2}**	0xF	Ex	{1}

5 Mem[0xB] = 5

3 InvReq 0xA

4 InvResp 0xA

6 ShResp 0xB, data=5

2 ShReq 0xB

Cache 0		
Tag	State	Data
0xA	**I**	3

Cache 1		
Tag	State	Data
0xF	M	1

Cache 2		
Tag	State	Data
0xB	**S**	**5**

| Core 0 |
|---|

Core 1

Core 2

1 LD 0xB

How many entries should the directory have?

Inexact Representations of Sharer Sets
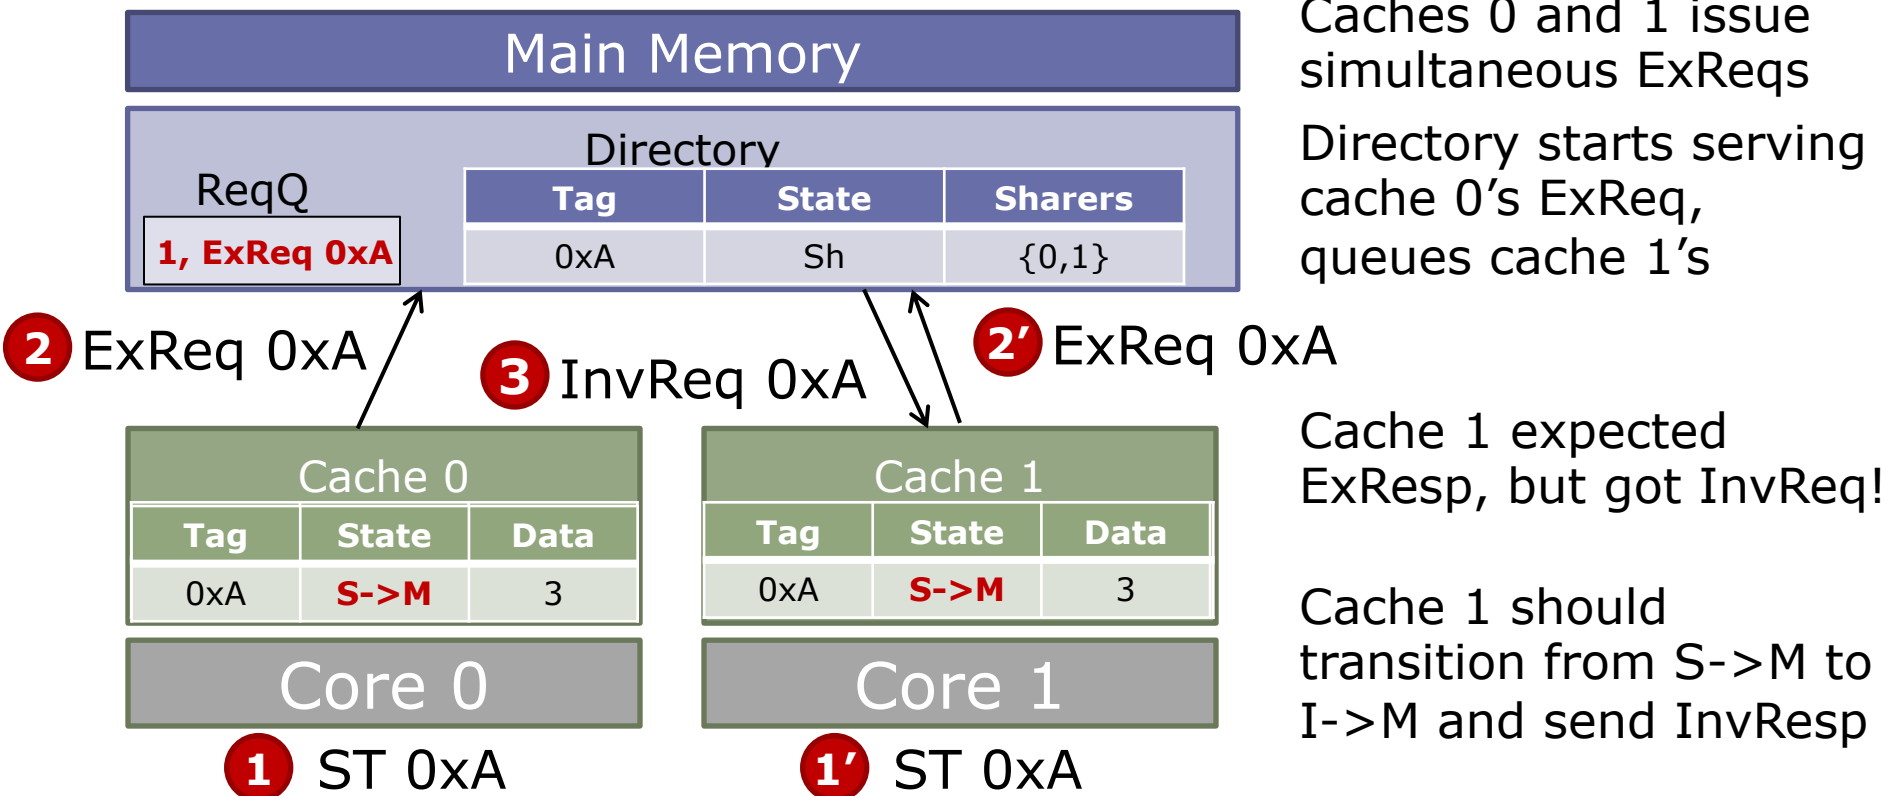
- Coarse-grain bit-vectors (e.g., 1 bit per 4 cores)

Sharer Set	0	0	0	0	0	0
	0-3	4-7	8-11	12-15	16-19	20-23

- Limited pointers: Maintain a few sharer pointers, on overflow mark 'all' and broadcast (or invalidate another sharer)

Sharer Set	0	8	14	33
	all	sharer 1	sharer 2	sharer 3

- Allow false positives (e.g., Bloom filters)

✓ Reduced area & energy
✗ Overheads still not scalable (these techniques simply play with constant factors)
✗ Inexact sharers → Broadcasts, invalidations or spurious invalidations and downgrades
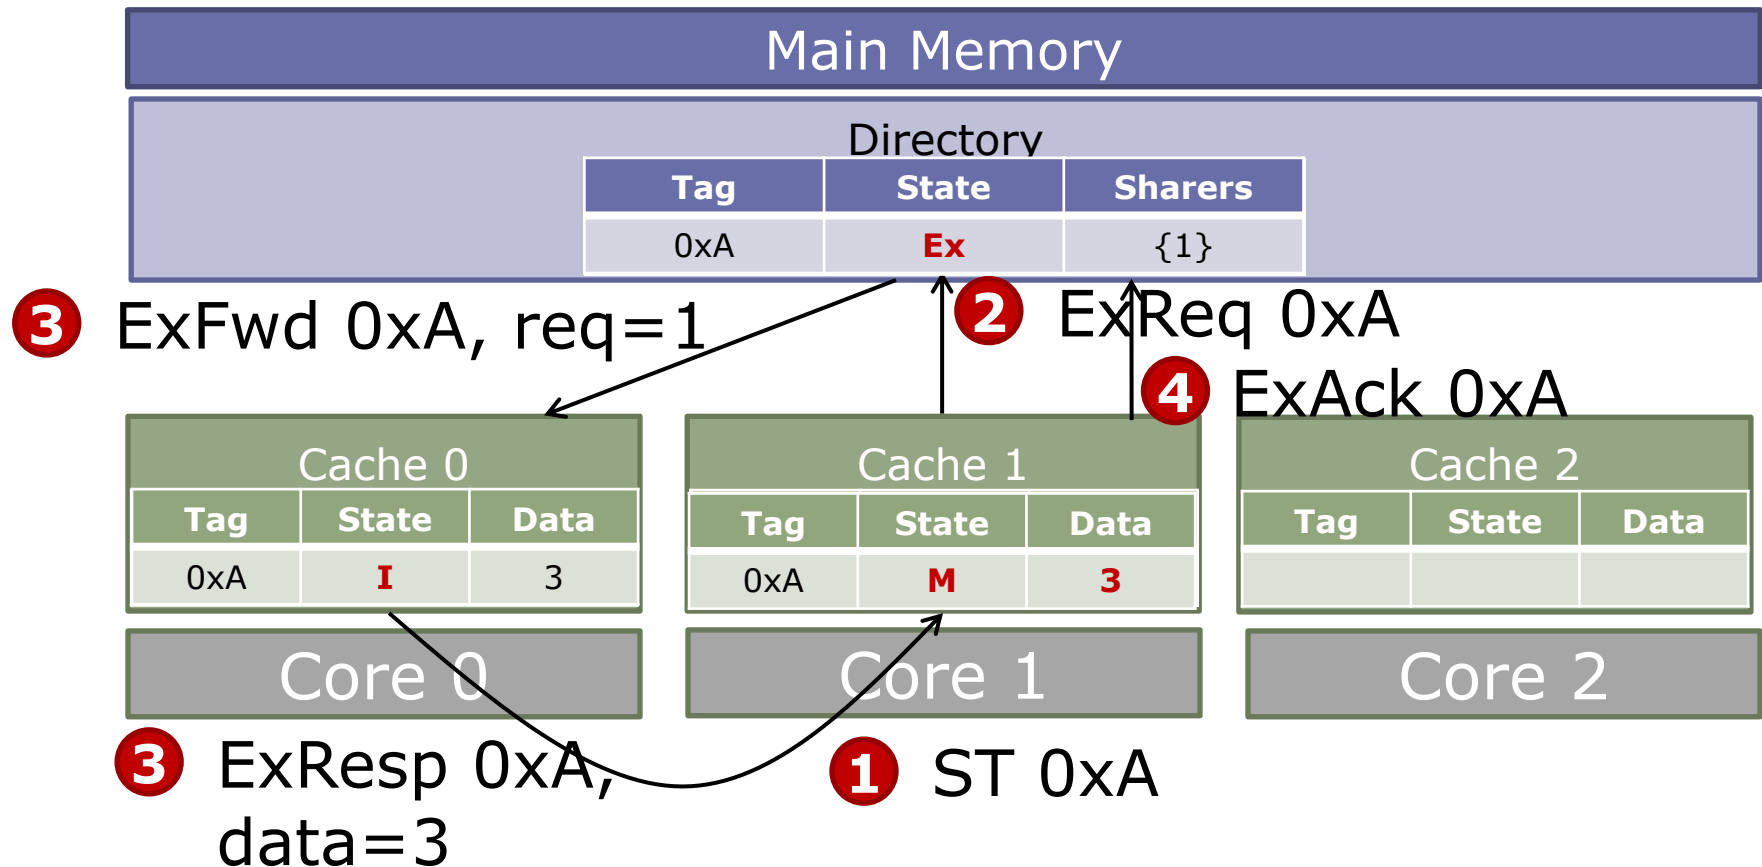
Protocol Races

- Directory serializes multiple requests for the same address
  - Same-address requests are queued or NACKed and retried
- But races still exist due to conflicting requests
- Example: Upgrade race

Caches 0 and 1 issue simultaneous ExReqs

Directory starts serving cache 0's ExReq, queues cache 1's

Cache 1 expected ExResp, but got InvReq!

Cache 1 should transition from S->M to I->M and send InvResp

Extra Hops and 3-Hop Protocols
Reducing Protocol Latency

- Problem: Data in another cache needs to pass through the directory, adding latency

- Optimization: Forward data to requester directly

Main Memory		

Directory

Tag	State	Sharers
0xA	Ex	{1}

3 ExFwd 0xA, req=1

2 ExReq 0xA

4 ExAck 0xA

Cache 0

Tag	State	Data
0xA	I	3

Cache 1

Tag	State	Data
0xA	M	3

Cache 2

Tag	State	Data

Core 0

Core 1

Core 2

3 ExResp 0xA, data=3

1 ST 0xA

Coherence in Multi-Level Hierarchies

- Can use the same or different protocols to keep coherence across multiple levels

- Key invariant: Ensure sufficient permissions in all intermediate levels

- Example: 8-socket Xeon E7 (8 cores/socket)

MESIF protocol
Snooping (QPI)

MESI protocol
L3 in-cache directory

In-Cache Directories

- Common multicore memory hierarchy:
  - 1+ levels of private caches
  - A shared last-level cache
  - Need to enforce coherence among private caches

- Idea: Embed the directory information in shared cache tags
  - Shared cache must be inclusive

Main Memory
Shared cache

Private cache	Private cache	...	Private cache
Core 0	Core 1		Core N

✓ Avoids tag overheads & separate lookups
✗ Can be inefficient if shared cache size >> sum(private cache sizes)
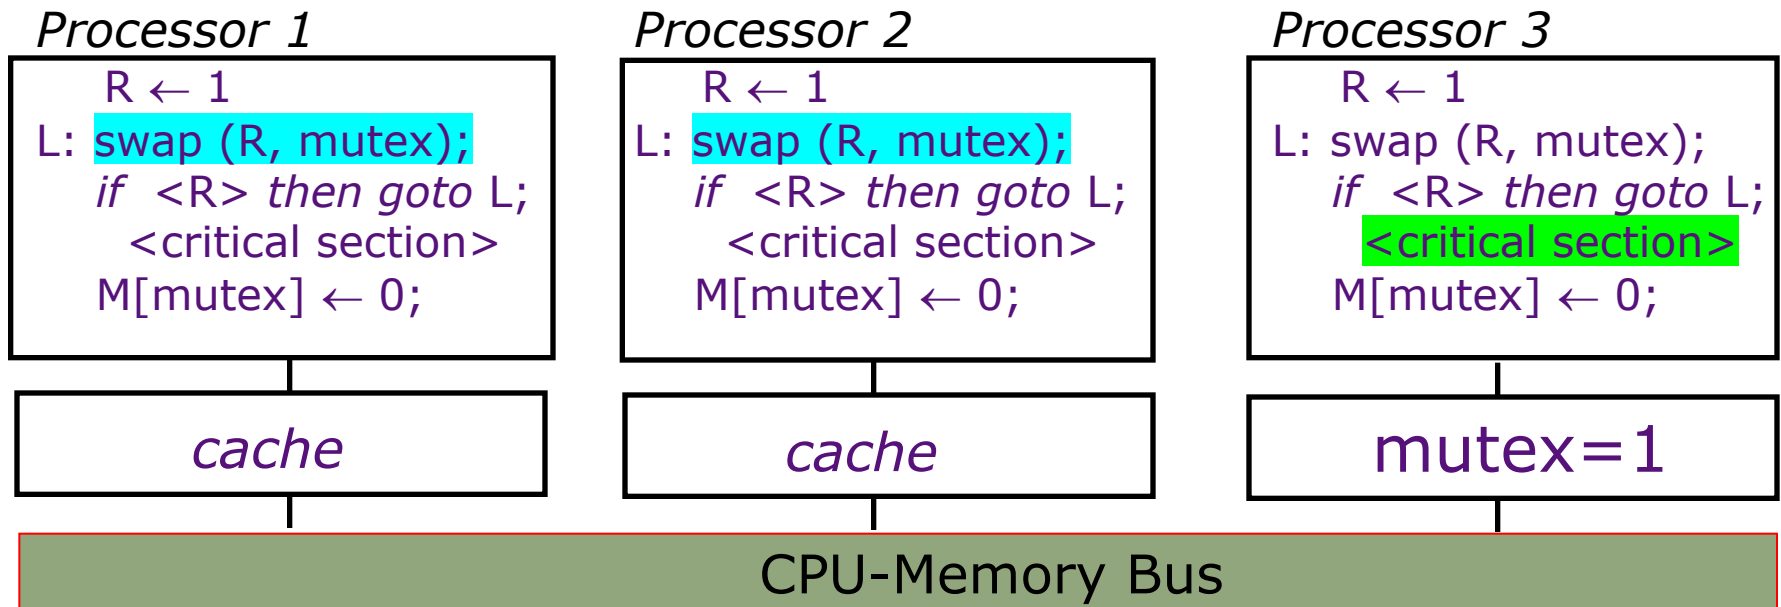
Avoiding Protocol Deadlock

- Protocols can cause deadlocks even if network is deadlock-free! (*more on this later*)

Example: Both nodes saturate all intermediate buffers with requests to each other, blocking responses from entering the network

- Solution: Separate *virtual networks*
  - Different sets of virtual channels and endpoint buffers
  - Same physical routers and links
  - More on this later (on-chip networks lectures)

- Most protocols require at least 2 virtual networks (for requests and replies), often >2 needed
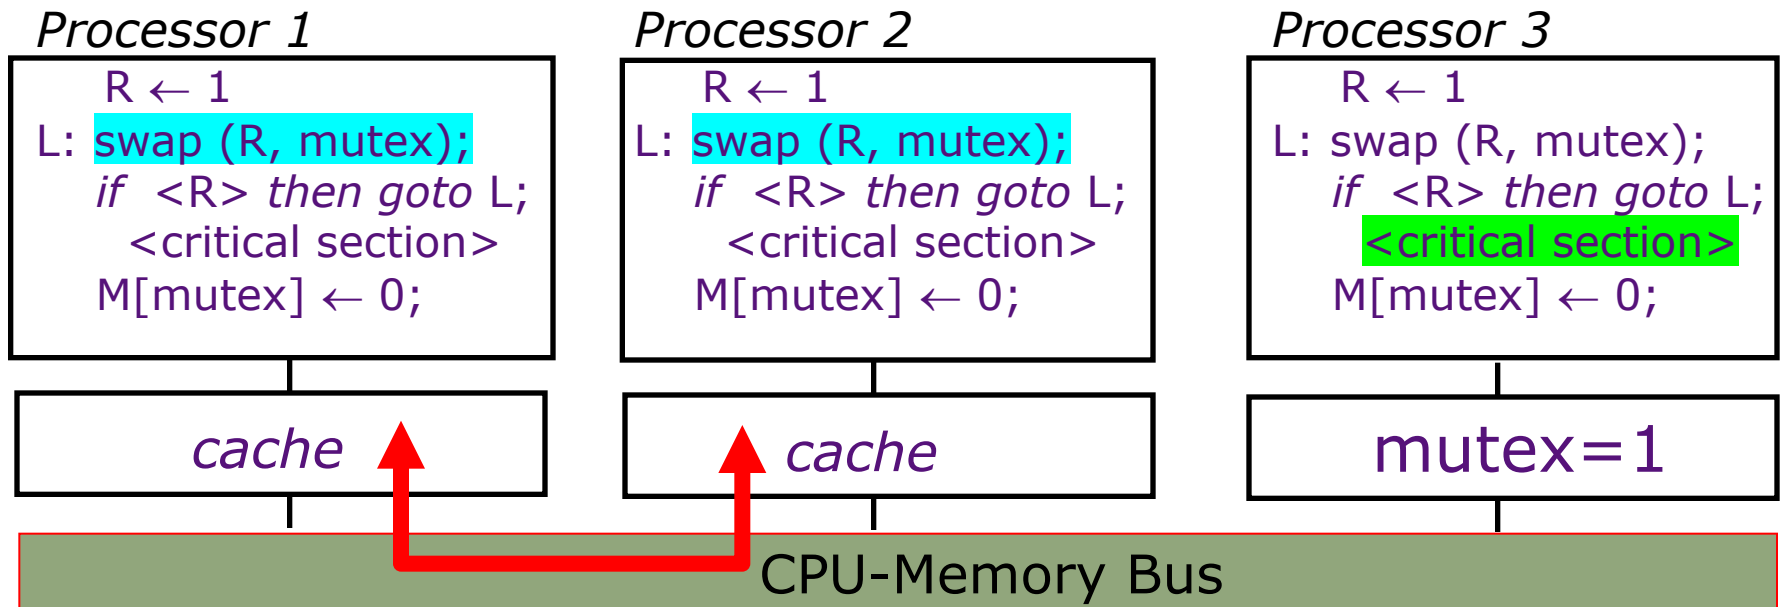
Coherence and Synchronization

Processor 1

```
     R ← 1
L: swap (R, mutex);
   if  <R> then goto L;
     <critical section>
   M[mutex] ← 0;
```

Processor 2

```
     R ← 1
L: swap (R, mutex);
   if  <R> then goto L;
     <critical section>
   M[mutex] ← 0;
```

Processor 3

```
     R ← 1
L: swap (R, mutex);
   if  <R> then goto L;
     <critical section>
   M[mutex] ← 0;
```

| cache | cache | mutex=1 |

CPU-Memory Bus

Implementing synchronization requires atomic read-modify-write operations, e.g., swap instruction

```
swap (R, mutex):
  load tmp, (mutex)
  store R, (mutex)
  R ← tmp
```
atomic

Our cache coherence protocol will introduce a performance issue here. What is the problem?

Coherence and Synchronization

Processor 1	Processor 2	Processor 3
R ← 1	R ← 1	R ← 1
L: swap (R, mutex);	L: swap (R, mutex);	L: swap (R, mutex);
if <R> then goto L;	*if <R> then goto* L;	*if <R> then goto* L;
<critical section>	<critical section>	<critical section>
M[mutex] ← 0;	M[mutex] ← 0;	M[mutex] ← 0;
cache	*cache*	mutex=1

CPU-Memory Bus

Cache coherence protocols will cause mutex to *ping-pong* between P1's and P2's caches.

Ping-ponging can be reduced by first reading the mutex location *(non-atomically)* and executing a swap only if it is found to be zero (test&test&set).

Implementing Atomic Instructions

- In general, an *atomic read-modify-write* instruction requires two memory (bus) operations without intervening memory operations by other processors

- Implementation options:
  - *With snoopy coherence, lock the bus → expensive*
  - *With directory-based coherence, lock the line in the cache (prevent invalidations or evictions until atomic op finishes) → complex*

- modern processors use
  load-reserve
  store-conditional

Load-reserve & Store-conditional

Special register(s) to hold reservation flag and address, and the outcome of store-conditional

Load-reserve R, (a):
    <flag, adr> ← <1, a>;
    R ← M[a];

Store-conditional (a), R:
    if <flag, adr> == <1, a>
    then cancel other procs'
                reservation on a;
            M[a] ← <R>;
            status ← succeed;
    else status ← fail;

If the cache receives an invalidation to the address in the reserve register, the reserve bit is set to 0
- Several processors may reserve 'a' simultaneously
- These instructions are like ordinary loads and stores with respect to the bus traffic

Load-Reserve/Store-Conditional

Swap implemented with Ld-Reserve/St-Conditional

```
#       Swap(R1, mutex):

L:      Ld-Reserve R2, (mutex)
        St-Conditional (mutex), R1
        if (status == fail) goto L
        R1 <- R2
```

Performance:
Load-reserve & Store-conditional

The total number of coherence transactions is not necessarily reduced, but splitting an atomic instruction into load-reserve & store-conditional:

- *increases utilization* (and reduces processor stall time), especially in split-transaction buses and directories

- *reduces cache ping-pong effect* because processors trying to acquire a semaphore do not have to perform stores each time

Thank you!

Next Lecture:
Consistency and
Relaxed Memory Models